

Internet Architecture: Design Philosophy - Then and Now

Partial credit goes to Prof. Z. Zhang

Internet Design Philosophy

1

Internet Philosophy and Design Principles

Architecture: the big picture

Goals:

- identify, study principles that can guide network architecture
- "bigger" issues than specific protocols or implementation tricks
- *synthesis*: the really big picture

2


Key questions

- How to decompose the complex system functionality into protocol layers?
- Which functions placed *where* in network, at which layers?
- Can a function be placed at multiple levels ?


Answer these questions in context of
Internet, telephone network

3

Common View of the Telco Network



Common View of the IP Network



5

Readings: Saltzer84

- End-to-end argument
 - Better to implement functions close to application
 - ... except when performance requires otherwise
- Why?
 - ...
- What should be the "end" for network "functionalities", e.g., routing?
 - Router?
 - End host?
 - Enterprise edge?
 - Autonomous System?

6


Internet End-to-End Argument

According to [Saltzer84]:

- "...functions placed at the lower levels may be *redundant* or of *little value* when compared to the cost of providing them at the lower level..."
- "...sometimes an *incomplete* version of the function provided by the communication system (lower levels) may be useful as a *performance enhancement*..."
- This leads to a philosophy diametrically opposite to the telephone world of dumb end-systems (the telephone) and intelligent networks.

7

Example: Reliable File Transfer



- Solution 1: make each step reliable, and then concatenate them
- Solution 2: each step unreliable: end-to-end check and retry

8

E2E Example: File Transfer

- Even if network guaranteed reliable delivery
 - Need to provide end-to-end checks
 - E.g., network card may malfunction
 - The receiver has to do the check anyway!
- Full functionality can only be entirely implemented at application layer; no need for reliability from lower layers
- Does FTP look like E2E file transfer?
 - TCP provides reliability between kernels not disks

Discussion

- Solution 1 not good enough!
 - what happens if the sender or/and receiver misbehave?
- so receiver has to do check anyway!
- Thus, full functionality can be entirely implemented at application layer; *no* need for reliability from lower layers

10

Discussion

Q: Is there any reason to implement reliability at lower layers?

A: Yes, but only to improve performance

- Example:
 - assume high error rate in network
 - reliable communication service at data link layer might help (why)?
 - fast detection /recovery of errors

11

E2E Argument: Interpretations

- One interpretation:
 - A function can only be completely and correctly implemented with the knowledge and help of the applications *standing at the communication endpoints*
- Another: (more precise...)
 - a system (or subsystem level) should consider only functions that can be *completely and correctly* implemented within it.
- Alternative interpretation: (also correct ...)
 - Think twice before implementing a functionality that you believe that is useful to an application at a lower layer
 - If the application can implement a functionality correctly, implement it a lower layer *only* as a performance enhancement

12

Internet & End-to-End Argument

- network layer provides one simple service: best effort datagram (packet) delivery
- transport layer at network edge (TCP) provides end-end error control
 - performance enhancement used by many applications (which could provide their own error control)
- all other functionalities ...
 - all application layer functionalities
 - network services: DNS implemented at application level

13

Internet & End-to-End Argument

Discussion: congestion control, "error" control, flow control: why at transport, rather than link or application layers?

- Claim: common functions should migrate down the stack
 - Everyone shares same implementation: no need to redo it (reduces bugs, less work, etc...)
 - Knowing everyone is doing the same thing, can help
- congestion control too important to leave up to application/user: true but hard to police
 - TCP is "outside" the network; compliance is "optional"
 - We do this for fairness (but realize that people could cheat)
- Why error control, flow control in TCP, not (just) in app

14

Trade-offs

- application has more information about the data and semantics of required service (e.g., can check only at the end of each data unit)
- lower layer has more information about constraints in data transmission (e.g., packet size, error rate)
- Note: these trade-offs are a direct result of layering!

15

End-to-End Argument: Critical Issues

- end-to-end principle emphasizes:
 - function placement
 - correctness, completeness
 - overall system costs
- Philosophy: if application can do it, don't do it at a lower layer -- application best knows what it needs
 - add functionality in lower layers iff (1) used by and improves performances of many applications, (2) does not hurt other applications
- allows cost-performance tradeoff

16

End-to-End Argument: Discussion

- end-end argument emphasizes correctness & completeness, but *not*
 - complexity: is complexity at edges result in a "simpler" architecture?
 - evolvability, ease of introduction of new functionality: ability to evolve because easier/cheaper to add new edge applications than change routers?
 - technology penetration: simple network layer makes it "easier" for IP to spread everywhere

17

Summary: End-to-End Arguments

- If the application can do it, don't do it at a lower layer -- anyway the application knows the best what it needs
 - add functionality in lower layers iff it is (1) used and improves performances of a large number of applications, and (2) does not hurt other applications
- Success story: Internet
 - But ...

18

Readings: Clark88

- Basic story of Clark88
 - Enumerate (and prioritize) system goals
 - ... and see what decisions that leads you to make
- Clark88 doesn't say much about routing, network trouble-shooting, security, etc., but
 - "Some of the most significant problems with the Internet today relate to lack of sufficient tools for distributed management, especially in the area of routing."
- What should be goals & priorities for routing, network trouble-shooting, security?
 - ...

19

Internet Design Philosophy (Clark'88)

In order of importance:

- 0 Connect existing networks
 - initially ARPANET and ARPA packet radio network
1. Survivability
 - ensure communication service even with network and router failures
2. Support multiple types of services
3. Must accommodate a variety of networks
4. Allow distributed management
5. Allow host attachment with a low level of effort
6. Be cost effective
7. Allow resource accountability

Different ordering of priorities would make a different architecture!

20

0. connect existing networks

Internet: virtualizing local networks

1974: multiple unconnected networks

- ARPAnet
- data-over-cable networks
- packet satellite network (Aloha)
- packet radio network

.. differing in:

- addressing conventions
 - packet formats
 - error recovery
 - routing
-

21

Challenge 1: Address Formats

- Map one address format to another?
 - Bad idea → many translations needed
 - Provide one common format
 - Map lower level addresses to common format
-


Challenge 2: Different Packet Sizes

- Define a maximum packet size over all networks?
 - Either inefficient or high threshold to support
 - Implement fragmentation/re-assembly
 - Who is doing fragmentation?
 - Who is doing re-assembly?
-

Gateway Alternatives

- Translation
 - Difficulty in dealing with different features supported by networks
 - Scales poorly with number of network types (N^2 conversions)
 - Standardization
 - "IP over everything" ([Design Principle 1](#))
 - Minimal assumptions about network
 - Hourglass design
-

Cerf & Kahn: Interconnecting two networks




- "...interconnection must preserve intact the internal operation of each network."
- "...the interface between networks must play a central role in the development of any network interconnection strategy. We give a special name to this interface that performs these functions and call it a GATEWAY."
- "...prefer that the interface be as simple and reliable as possible, and deal primarily with passing data between networks that use different packet-switching strategies"
- "...address formats is a problem between networks because the local network addresses of TCP's may vary substantially in format and size. A uniform internetwork TCP address space, understood by each GATEWAY and TCP, is essential to routing and delivery of internetwork packets."

25


Cerf & Kahn: Interconnecting two networks

- Internetwork layer:**
 - addressing: internetwork appears as a single, uniform entity, despite underlying local network heterogeneity
 - network of networks
- Gateway:**
 - "embed internetwork packets in local packet format or extract them"
 - route (at internetwork level) to next gateway



26

Historical Aside: Proposed Internetwork packet in 1974:



27

Cerf & Kahn's Internetwork Architecture

- two layers of addressing: internetwork and local network
- new layer makes everything homogeneous
- underlying local network technology (cable, satellite, 56K modem) is "invisible" at internetwork layer


28

IP Standardization

- Minimum set of assumptions for underlying net
 - Minimum packet size
 - Reasonable delivery odds, but not 100%
 - Some form of addressing unless point to point
- Important non-assumptions:
 - Perfect reliability
 - Broadcast, multicast
 - Priority handling of traffic
 - Internal knowledge of delays, speeds, failures, etc
- Also achieves Goal 3: Supporting Varieties of Networks

IP Hourglass


- Need to interconnect many existing networks
- Hide underlying technology from applications
- Decisions:
 - Network provides minimal functionality
 - "Narrow waist"



Tradeoff: No assumptions, no guarantees.

IP Layering (Principle 2)

- Relatively simple
- Sometimes taken too far



1. Survivability

- Continue to operate even in the presence of network failures (e.g., link and router failures)
 - as long as network is not partitioned, two endpoints should be able to communicate
 - any other failure (excepting network partition) should be transparent to endpoints
- Decision: maintain e-e transport state only at end-points
 - eliminate the problem of handling state inconsistency and performing state restoration when router fails
- Internet: stateless network architecture
 - No notion of a session/call at network layer
 - Grade: A-, because convergence times are relatively slow
 - BGP can take minutes to converge
 - IS-IS OSPF take ~ 10 seconds

Survivability

- If network disrupted and reconfigured
 - Communicating entities should not care!
 - No higher-level state reconfiguration
- How to achieve such reliability?
 - Where can communication state be stored?

	Network	Host
Failure handing	Replication	"Fate sharing"
Net Engineering	Tough	Simple
Switches	Maintain state	Stateless
Host trust	Less	More

Principle 3: Fate Sharing



- Lose state information for an entity if and only if the entity itself is lost.
- Examples:
 - OK to lose TCP state if one endpoint crashes
 - NOT okay to lose if an intermediate router reboots
 - Is this still true in today's network?
 - NATs and firewalls
- Survivability compromise: Heterogeneous network
 - less information available to end hosts and Internet level recovery mechanisms

Principle 4: Soft-state

- Soft-state
 - Announce state
 - Refresh state
 - Timeout state
- Penalty for timeout - poor performance
- Robust way to identify communication flows
 - Possible mechanism to provide non-best effort service
- Helps survivability

2. Types of Service

- **Principle 6:** network layer provides one simple service: best effort datagram (packet) delivery
 - All packets are treated the same
- Relatively simple core network elements
- Building block from which other services (such as reliable data stream) can be built
- Contributes to scalability of network
- No QoS support assumed from below
 - In fact, some underlying nets only supported reliable delivery
 - Made Internet datagram service less useful!
 - Hard to implement without network support
 - QoS is an ongoing debate...

ToS: TCP vs. UDP

- Original Internet model: "TCP/IP" one layer
- add UDP to TCP to better support other apps
 - e.g., "real-time" applications
- arguably main reason for separating TCP, IP
- datagram abstraction: lower common denominator on which other services can be built
 - service differentiation was considered (remember ToS?), but this has never happened on the large scale (Why?)
- proven to allow lots of applications to be invented and flourish (except MM, but maybe that's not a transport service issue)

37

3. Variety of Networks

- Very successful (why?)
 - because the minimalist service; it requires from underlying network only to deliver a packet with a "reasonable" probability of success
 - ...does not require:
 - reliability
 - in-order delivery
 - The mantra: IP over everything
 - Then: ARPANET, X.25, DARPA satellite network..
 - Now: ATM, SONET, WDM...
- A: can't name a link layer technology that IP doesn't run over (carrier pigeon RFC)

38

Other Goals

- Allow distributed management
 - Administrative autonomy: IP interconnects networks
 - each network can be managed by a different organization
 - different organizations need to interact only at the boundaries
 - ... but this model complicates routing
 - A for implementation, B for concept (disagreement)
- Cost effective
 - sources of inefficiency
 - header overhead
 - retransmissions
 - routing
 - ...but "optimal" performance never been top priority
 - A

39

Other Goals (Cont)

- Low cost of attaching a new host
 - not a strong point → higher than other architecture because the intelligence is in hosts (e.g., telephone vs. computer)
 - bad implementations or malicious users can produce considerably harm (remember fate-sharing?)
 - C: but things are improving with DHCP, auto-configurations. Looks like a higher grade in future
- Accountability
 - Internet gets an "F"


40

Summary: Minimalist Approach

- **Dumb network**
 - IP provide minimal functionalities to support connectivity
 - Addressing, forwarding, routing
- **Smart end system**
 - Transport layer or application performs more sophisticated functionalities
 - Flow control, error control, congestion control
- **Advantages**
 - Accommodate heterogeneous technologies (Ethernet, modem, satellite, wireless)
 - Support diverse applications (telnet, ftp, Web, X windows)
 - Decentralized network administration

Summary: Internet Architecture

- packet-switched datagram network
- IP is the glue (network layer overlay)
- IP hourglass architecture
 - all hosts and routers run IP
- stateless architecture
 - no per flow state inside network



42

Summary: Minimalist Approach

- **Dumb network**
 - IP provide minimal functionalities to support connectivity
 - addressing, forwarding, routing
- **Smart end system**
 - transport layer or application performs more sophisticated functionalities
 - flow control, error control, congestion control
- **Advantages**
 - accommodate heterogeneous technologies (Ethernet, modem, satellite, wireless)
 - support diverse applications (telnet, ftp, Web, X windows)
 - decentralized network administration

43

But that was **yesterday**

..... what about **tomorrow?**

44

What About the Future?

- Datagram not the best abstraction for:
 - resource management, accountability, QoS
- new abstraction: **flow** (see IPv6)
 - but no one knows what a flow is
- routers require to maintain per-flow state
- state management: recovering lost state is hard
- here we see the first proposal of "soft state"!
 - **soft-state**: end-hosts responsible to maintain the state

45

Rethinking Internet Design

What's changed?

- **operation in untrustworthy world**
 - endpoints can be malicious
 - If endpoint not trustworthy, but want trustworthy network -> more mechanism in network core
 - Trust and security a big issue today!
- **more demanding applications**
 - end-end best effort service not enough
 - new service models in network (Intserv, Diffserv)?
 - new application-level service architecture built on top of network core (e.g., CDN, p2p)?
 - wireless and mobility

46

Rethinking Internet Design ...

What's changed?

- **ISP service differentiation**
 - ISP doing more (than other ISPs) in core is competitive advantage
- **rise of third party involvement**
 - interposed between endpoints (even against will)
 - e.g., Chinese government, US recording industry
- **new technologies (wireless, optical ...)**
- **limited capability devices (e.g., PDA, smart phones, sensors,), or perhaps also less "sophisticated" users**

All these changes motivate shift away from end-end!

47

What's at stake?

"At issue is the conventional understanding of the 'Internet philosophy'

- freedom of action
- user empowerment
- end-user responsibility for actions taken
- lack of control "in" the net that limit or regulate what users can do

The end-end argument fostered that philosophy because they enable the freedom to innovate, install new software at will, and run applications of the users choice"

[Blumenthal and Clark, 2001]

48

Technical response to changes

- Add functions to the network core ("middleboxes"):
 - filtering firewalls
 - application-level firewalls, web caches and proxies
 - NAT boxes
 - active networking
 - ...
 - Add "infrastructure services"
 - e.g., DNS,
 - (application-specific) content distribution networks (CDNs)
- ... All operate within network, making use of application-level information
- which addresses can do what at application level?
 - If addresses have meaning to applications, NAT must "understand" that meaning

49

Next Week

- Review lecture notes and the required readings for this week:
 - [Saltzer84] and [Clark88]
 - also [Clark:Tussle] and [CerfKahn] if you have time
- Questions for you to think about:
 - What are the "architectural" advantages of Internet, and also its limitations?
 - If you can redesign it, how would you do it?

50