



# Advanced Computer Networks

---

QoS



# Overview

---

- Why QOS?
- Integrated services
- RSVP
- Adaptive applications
- Differentiated services



# Motivation

---

- Internet currently provides one single class of **“best-effort” service**
  - No assurances about delivery
- Existing applications are *elastic*
  - Tolerate delays and losses
  - Can adapt to congestion
- Future “real-time” applications may be *inelastic*



# Inelastic Applications

---

- Continuous media applications
  - **Lower and upper limit** on acceptable performance.
  - BW below which video and audio are not intelligible
  - Internet telephones, teleconferencing with high delay (200 - 300ms) impair human interaction
- Hard real-time applications
  - Require **hard limits on performance**
  - E.g. control applications






# Why a New Service Model?

---


- What is the **basic objective** of network design?
  - Maximize total bandwidth? Minimize latency?
  - **Maximize user satisfaction** – the total **utility** given to users
- What does utility vs. bandwidth look like?
  - Must be non-decreasing function
  - Shape depends on application

# Utility Curve Shapes



Stay to the right and you  
are fine for all curves


# Utility curve – Elastic traffic



**Does equal allocation of bandwidth maximize total utility?**



# Admission Control

- If  $U(\text{bandwidth})$  is concave  
→ elastic applications
    - Incremental utility is decreasing with increasing bandwidth
    - Is always advantageous to have more flows with lower bandwidth
      - No need of admission control;
- This is why the Internet works!





# Utility Curves – Inelastic traffic



**Does equal allocation of  
bandwidth maximize total utility?**




# Why a New Service Model?

---

- Given the shape of different utility curves – clearly equal allocation of bandwidth does not maximize total utility
- In fact, desirable rate for some flow may be 0.

# Admission Control

- If  $U$  is convex  $\rightarrow$  inelastic applications
  - $U$ (number of flows) is no longer monotonically increasing
  - Need admission control to maximize total utility
- **Admission control**  $\rightarrow$  deciding when the addition of new people would result in reduction of utility
  - Basically avoids overload





# Admission Control

---

## ■ Caveats

- Admission control can only turn away new requests → sometimes it may be better to terminate an existing flow
- $U(0) \neq 0 \rightarrow$  users tend to be very unhappy with no service – perhaps  $U$  should be discontinuous here

## ■ Alternative → overprovision the network

- Problem: high variability in usage patterns
- “Leading-edge” users make it costly to overprovision
- Having admission control seems to be a better alternative



# Other QOS principles

---

1. Admission Control
2. Marking of packets is needed to distinguish between different classes.
3. Protection (isolation) for one class from another.
4. While providing isolation, it is desirable to use resources as efficiently as possible  
→ sharing.



# Overview

---

- Why QOS?
- Integrated services
- Adaptive applications
- Differentiated services



# Components of Integrated Services

---

1. **Type of commitment**

**What does the network promise?**

2. Packet scheduling

How does the network meet promises?

3. Service interface

How does the application describe what it wants?

4. Establishing the guarantee

How is the promise communicated to/from the network


How is admission of new applications controlled?



# 1. Type of commitment

---

What kind of promises/services should network offer?



Depends on the **characteristics of the applications** that will use the network ....





# Playback Applications

---

- Sample signal → packetize → transmit → buffer → playback
  - Fits most multimedia applications
- Performance concern:
  - Jitter – variation in end-to-end delay
    - Delay = fixed + variable = (propagation + packetization) + queuing
- Solution:
  - Playback point – delay introduced by buffer to hide network jitter



# Characteristics of Playback Applications

---

- In general lower delay is preferable.
- Doesn't matter when packet arrives as long as it is before playback point
- Network guarantees (e.g. bound on jitter) would make it easier to set playback point
- Applications can tolerate some loss



# Applications Variations

---

- Rigid & adaptive applications
  - Rigid – set fixed playback point
  - Adaptive – adapt playback point
    - Gamble that network conditions will be the same as in the past
    - Are prepared to deal with errors in their estimate
    - Will have an earlier playback point than rigid applications
- Tolerant & intolerant applications
  - Tolerance to brief interruptions in service
- 4 combinations



# Applications Variations

---

## **Really only two classes of applications**

- 1) Intolerant and rigid**
- 2) Tolerant and adaptive**

Other combinations make little sense

- 3) Intolerant and adaptive
  - Cannot adapt without interruption
- 4) Tolerant and rigid
  - Missed opportunity to improve delay

**So what service classes should the network offer?**



# Type of Commitments

---

- **Guaranteed** service
  - For **intolerant and rigid** applications
  - Fixed guarantee, network meets commitment as long as clients send at match traffic agreement
- **Predicted** service
  - For **tolerant and adaptive** applications
  - Two components
    - If conditions do not change, commit to current service
    - If conditions change, take steps to deliver consistent performance (help apps minimize playback delay)
    - Implicit assumption – network does not change much over time
- **Datagram/best effort service**



# Components of Integrated Services

---

1. Type of commitment  
What does the network promise?
2. **Packet scheduling**  
**How does the network meet promises?**
3. **Service interface**  
**How does the application describe what it wants?**
4. Establishing the guarantee  
How is the promise communicated to/from the network  
How is admission of new applications controlled?




# Scheduling for Guaranteed Traffic

---

- Use **token bucket filter** to characterize traffic
  - Described by rate  $r$  and bucket depth  $b$
- Use **WFQ** at the routers
- Parekh's bound for worst case queuing delay =  $b/r$

# Token Bucket Filter



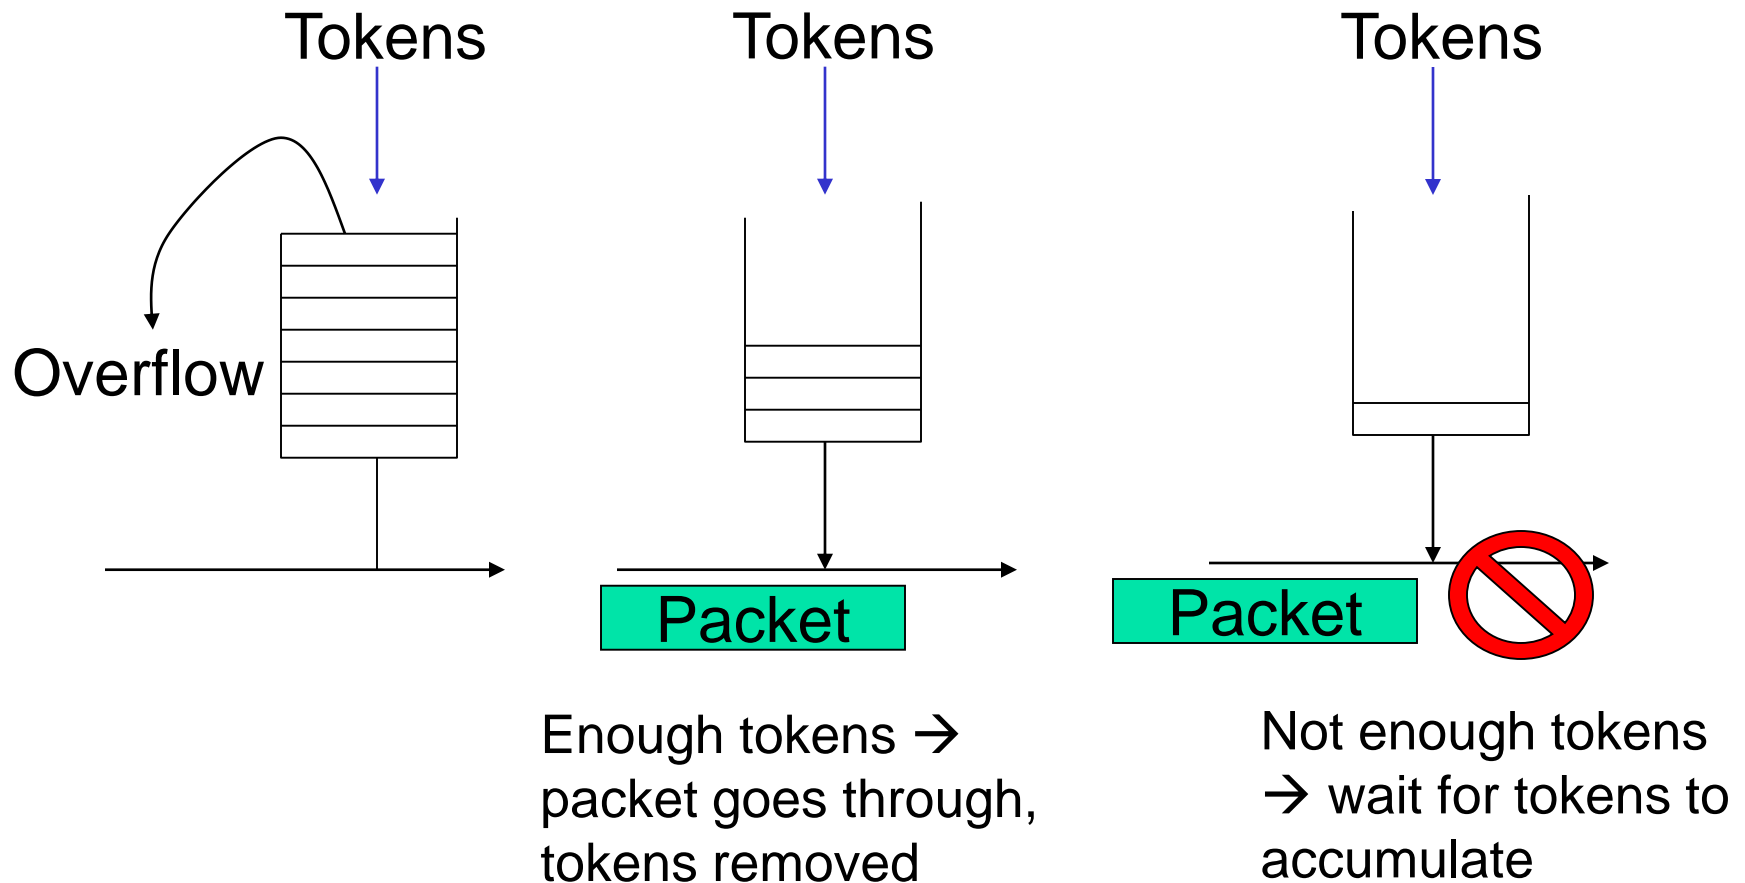
## Operation:

- If bucket fills, tokens are discarded
- Sending a packet of size  $P$  uses  $P$  tokens
- If bucket has  $P$  tokens, packet sent at max rate, else must wait for tokens to accumulate

Parekh's bound for worst case queuing delay =  $b/r$



# Token Bucket Operation





# Guarantee Proven by Parekh

---

- Given:
  - Flow  $i$  shaped with token bucket and leaky bucket rate control (depth  $b$  and rate  $r$ )
  - Network nodes do WFQ
- Cumulative queuing delay  $D_i$  suffered by flow  $i$  has upper bound
  - **$D_i < b/r$** , (where  $r$  may be much larger than average rate)
  - Assumes that  $\forall r < \text{link speed at any router}$
  - All sources limiting themselves to  $r$  will result in no network queuing




# Token Bucket Characteristics

---

- On the long run, rate is limited to  $r$
- On the short run, a burst of size  $b$  can be sent
- Amount of traffic entering at interval  $T$  is bounded by:
  - Traffic =  $b + r * T$
- Information useful to admission algorithm

# Token Bucket Specs



Flow A:  $r = 1$  MBps,  $B = 1$  byte

Flow B:  $r = 1$  MBps,  $B = 1$  MB



# Unified Scheduling

---

- Assume 3 types of traffic: guaranteed, predictive, best-effort
- Scheduling: use WFQ in routers
- Each guaranteed flow gets its own queue
- All predicted service flows and best effort aggregates in single separate queue
  - Predictive traffic classes
    - Multiple FIFO+ queues
    - Worst case delay for classes separated by order of magnitude
    - When high priority needs extra bandwidth – steals it from lower class
  - Best effort traffic acts as lowest priority class



# Service Interfaces

---

## ■ Guaranteed Traffic

- Host specifies rate to network
- Why not bucket size  $b$ ?
  - If delay not good, ask for higher rate

## ■ Predicted Traffic

- Specifies  $(r, b)$  token bucket parameters
- Specifies delay  $D$  and loss rate  $L$
- Network assigns priority class
- Policing at edges to drop or tag packets
  - Needed to provide isolation – why is this not done for guaranteed traffic?
    - WFQ provides this for guaranteed traffic



## How to Choose Service – Implicit

---

Network could examine packets and **implicitly** determine service class

- No changes to end hosts/applications
- Fixed set of applications supported at any time
- Can't support applications in different uses/modes easily
- Violates layering/modularity



# How to Choose Service – Explicit

---

Applications could **explicitly** request service level

- Why would an application request lower service?
  - Pricing
  - Informal social conventions
  - Problem exists in best-effort as well → congestion control
- Applications must know network service choices
  - Difficult to change over time
  - All parts of network must support this → places greater burden on portability of IP





# Overview

---

- Why QOS?
- Integrated services
- RSVP
- Adaptive applications
- Differentiated services



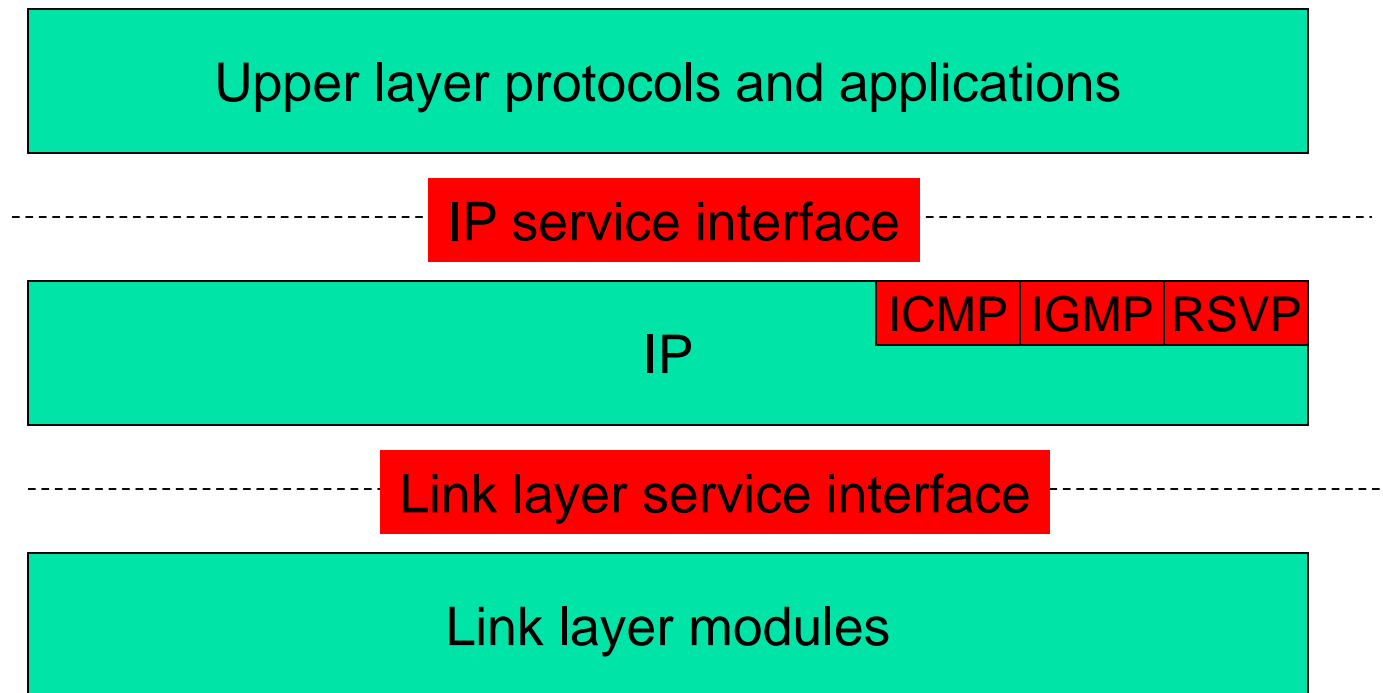
# Components of Integrated Services

---

1. Type of commitment  
What does the network promise?
2. Packet scheduling  
How does the network meet promises?
3. Service interface  
How does the application describe what it wants?
4. **Establishing the guarantee**  
**How is the promise communicated**  
**How is admission of new applications controlled?**



# Reservation Protocol: RSVP





## Role of RSVP

---

- Rides on top of unicast/multicast routing protocols
- Carries resource requests all the way through the network
- At each hop consults admission control and sets up reservation. Informs requester if failure



# RSVP Goals

---

- Used on connectionless networks
  - Should not replicate routing functionality
  - Should co-exist with route changes
- Support for multicast
  - Different receivers have different capabilities and want different QOS
  - Changes in group membership should not be expensive
  - Reservations should be aggregate – I.e. each receiver in group should not have to reserve
  - Should be able to switch allocated resource to different senders
- Modular design – should be generic “signaling” protocol
- Result
  - Receiver-oriented
  - Soft-state



# RSVP Service Model

---

- Make reservations for simplex data streams
- Receiver decides whether to make reservation
- Control msgs in IP datagrams (proto #46)
- PATH/RESV sent periodically to refresh soft state
- One pass:
  - Failed requests return error messages - receiver must try again
  - No e2e ack for success



# PATH Messages

---

- PATH messages carry sender's Tspec
  - Token bucket parameters
- Routers note the direction PATH messages arrived and set up *reverse path* to sender
- Receivers send RESV messages that follow reverse path and setup reservations
- If reservation cannot be made, user gets an error




# RESV Messages

---

- Forwarded via reverse path of PATH
- Queuing delay and bandwidth requirements
- Source traffic characteristics (from PATH)
- Filter specification
  - Which transmissions can use the reserved resources
- Router performs admission control and reserves resources
  - If request rejected, send error message



# PATH and RESV Messages





# Routing Changes

---

- Routing protocol makes routing changes
- In absence of route or membership changes, periodic PATH and RESV msgs refresh established reservation state
- When change, new PATH msgs follow new path, new RESV msgs set reservation
- Non-refreshed state times out automatically



# Overview

---

- Why QOS?
- Integrated services
- RSVP
- Adaptive applications
- Differentiated services



# Internet Video Today

---

- Client-server streaming
  - Skype video conferencing
  - Hulu
- DVD transfer
  - BitTorrent → P2P lecture
- Synchronized video (IPTV)
  - Overlay multicast → multicast lecture

# Client-Server Streaming: Adaptation Quality to Link



Long Time Scale

Short Time Scale

Content Negotiation




Server Selection



Adaptive Media

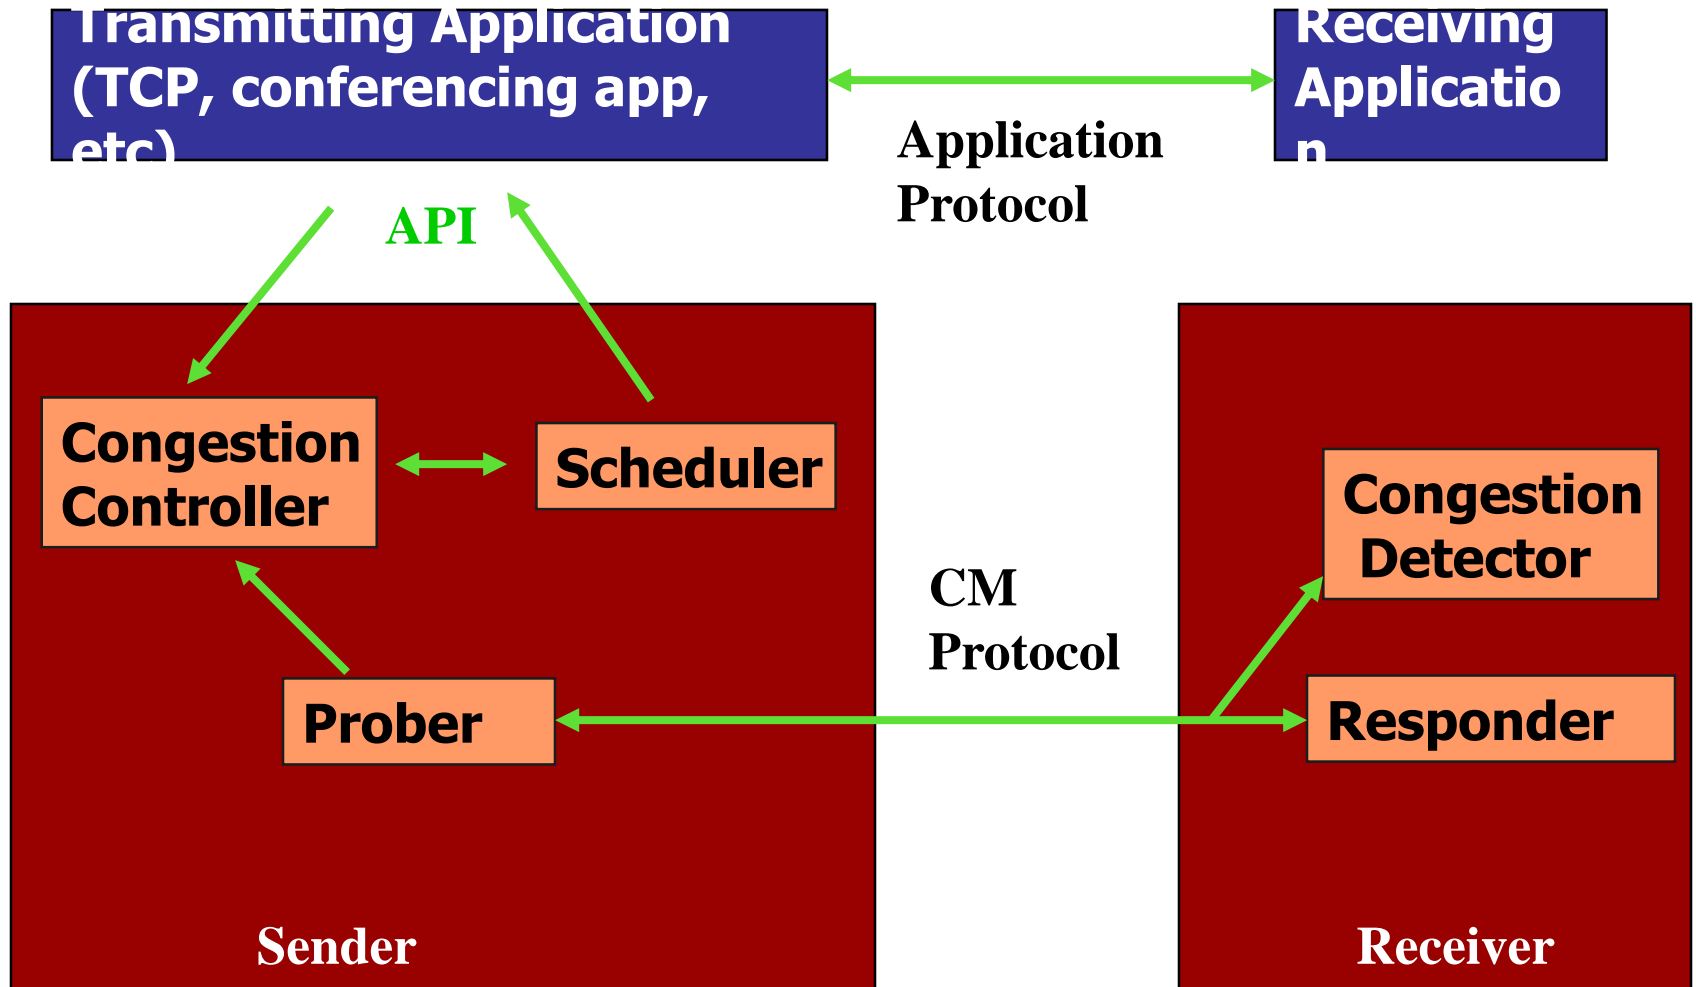
# Problems Adapting to Network State



- TCP hides network state
- New applications may not use TCP
  - Often do not adapt to congestion


**Need system that helps applications learn and adapt to congestion**

# Congestion Manager Architecture



# Transmission API

- Buffered send
  - `cm_send(data, length)`
- Request/callback-based send







# Transmission API (cont.)

---

- Request API: asynchronous sources

```
wait for (some_events) {  
    get_data( );  
    send( );  
}
```

- Synchronous sources

```
do_every_t_ms {  
    get_data( );  
    send( );  
}
```

- Solution: `cmapp_update(rate, srtt)` callback



# Feedback about Network State

---

- Monitoring successes and losses
  - Application hints
  - Probing system
- Notification API (application hints)
  - Application calls `cm_update(nsent, nrecd, congestion indicator, rtt)`



# Overview

---

- Why QOS?
- Integrated services
- Adaptive applications
- Differentiated services

- Analogy:
  - Airline service, first class, coach, various restrictions on coach as a function of payment
- Best-effort expected to make up bulk of traffic, but revenue from first class important to economic base (will pay for more plentiful bandwidth overall)
- Not motivated by real-time! Motivated by economics and assurances



# Basic Architecture

---

- Agreements/service provided within a domain
  - Service Level Agreement (SLA) with ISP
- Edge routers do traffic conditioning
  - Perform per aggregate shaping and policing
  - Mark packets with a small number of bits; each bit encoding represents a class or subclass
- Core routers
  - Process packets based on packet marking and defined per hop behavior
- More scalable than IntServ
  - No per flow state or signaling



## Per-hop Behaviors (PHBs)

---

- Define behavior of individual routers rather than end-to-end services – there may be many more services than behaviors
- Multiple behaviors – need more than one bit in the header
- Six bits from IP TOS field are taken for Diffserv code points (DSCP)



# Per-hop Behaviors (PHBs)

---

- Two PHBs defined so far
- Expedited forwarding aka premium service (type P)
  - Possible service: providing a virtual wire
  - Admitted based on peak rate
  - Unused premium goes to best effort
- Assured forwarding (type A)
  - Possible service: strong assurance for traffic within profile & allow source to exceed profile
  - Based on expected capacity usage profiles
  - Traffic unlikely to be dropped if user maintains profile
  - Out-of-profile traffic marked



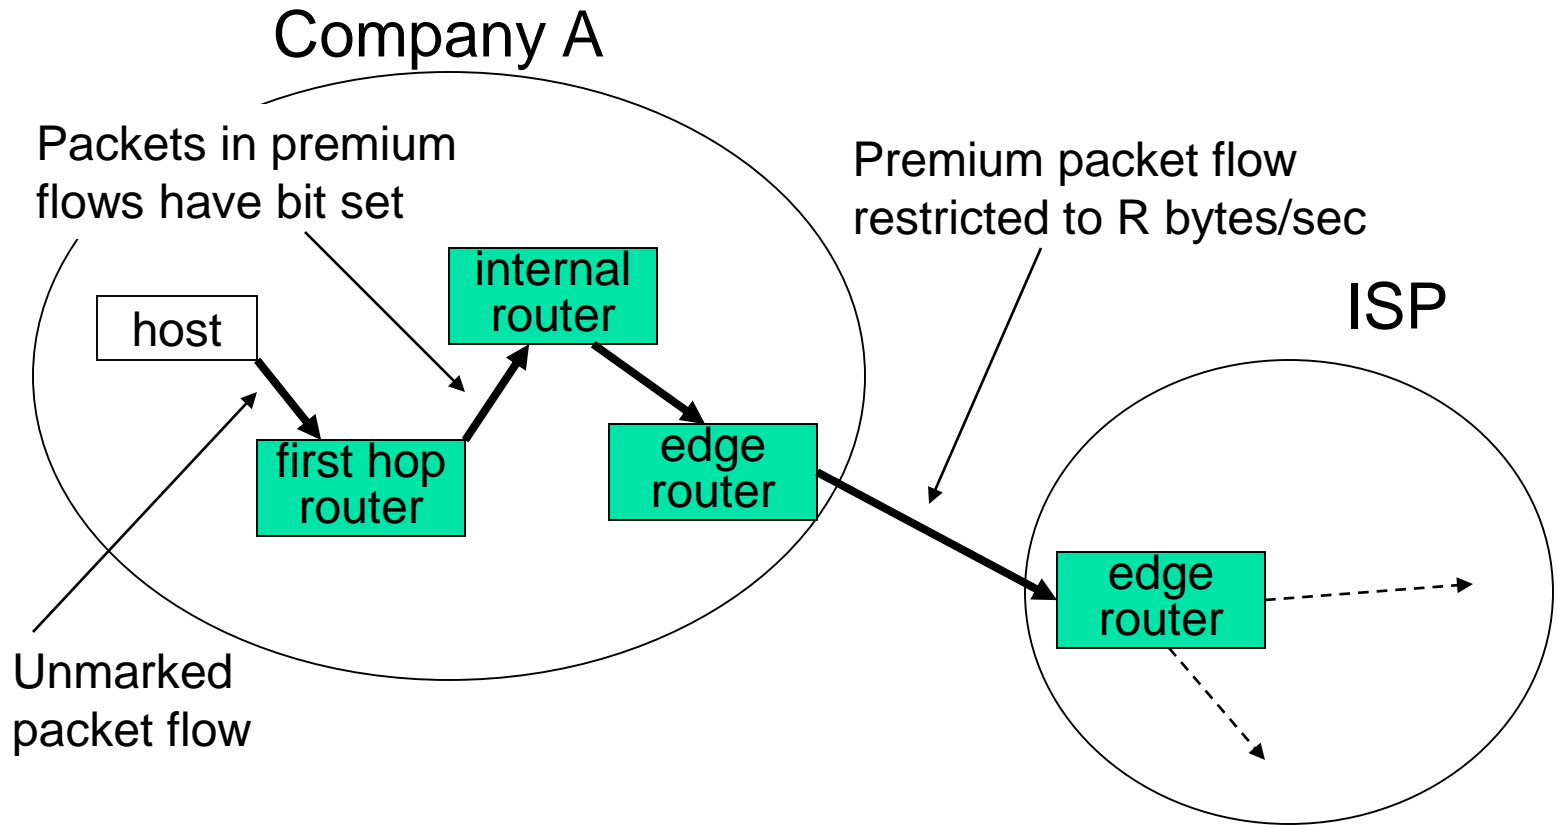
# Expedited Forwarding PHB

---

- User sends within profile & network commits to delivery with requested profile
  - Signaling, admission control may get more elaborate in future
- Rate limiting of EF packets at edges only, using token bucket to shape transmission
- Simple forwarding: classify packet in one of two queues, use priority
  - EF packets are forwarded with minimal delay and loss (up to the capacity of the router)



# Expedited Forwarding Traffic Flow





# Assured Forwarding PHB

---

- User and network agree to some traffic profile
  - Edges mark packets up to allowed rate as “in-profile” or low drop precedence
  - Other packets are marked with one of 2 higher drop precedence values
- A congested DS node tries to protect packets with a lower drop precedence value from being lost by preferably discarding packets with a higher drop precedence value
  - Implemented using RED with In/Out bit

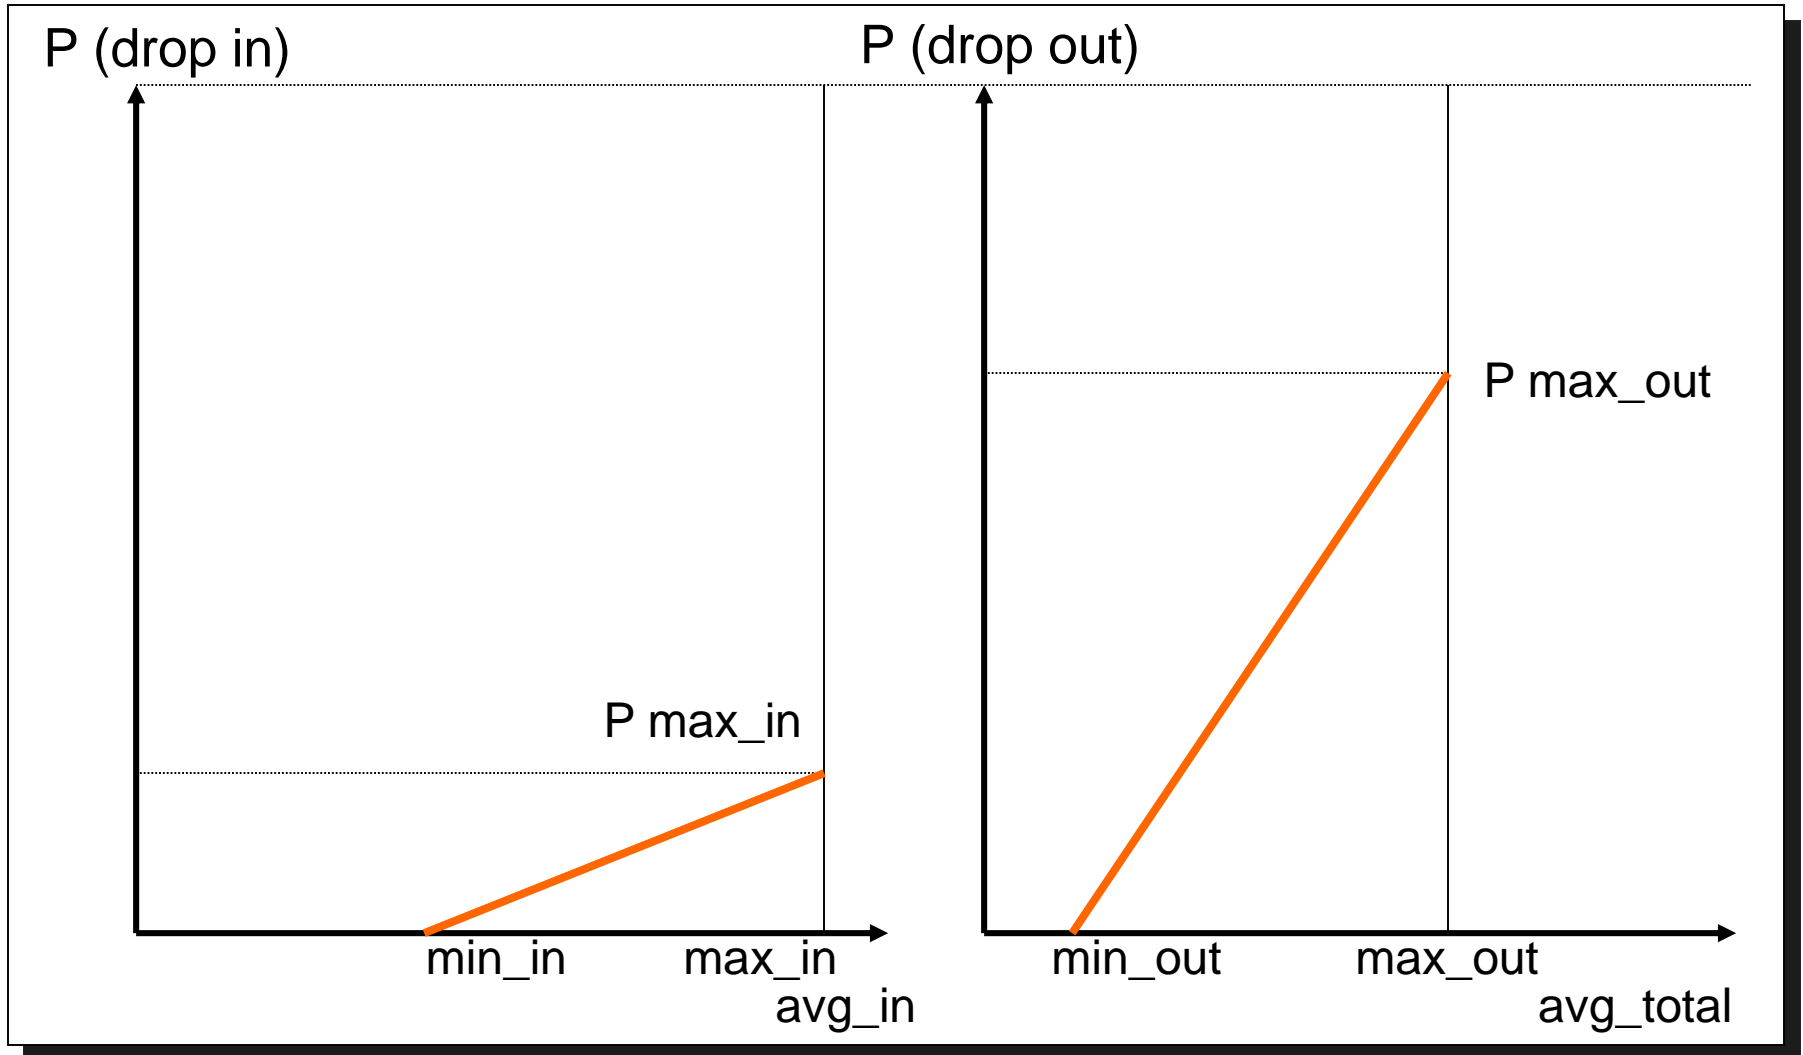


## Red with In or Out (RIO)

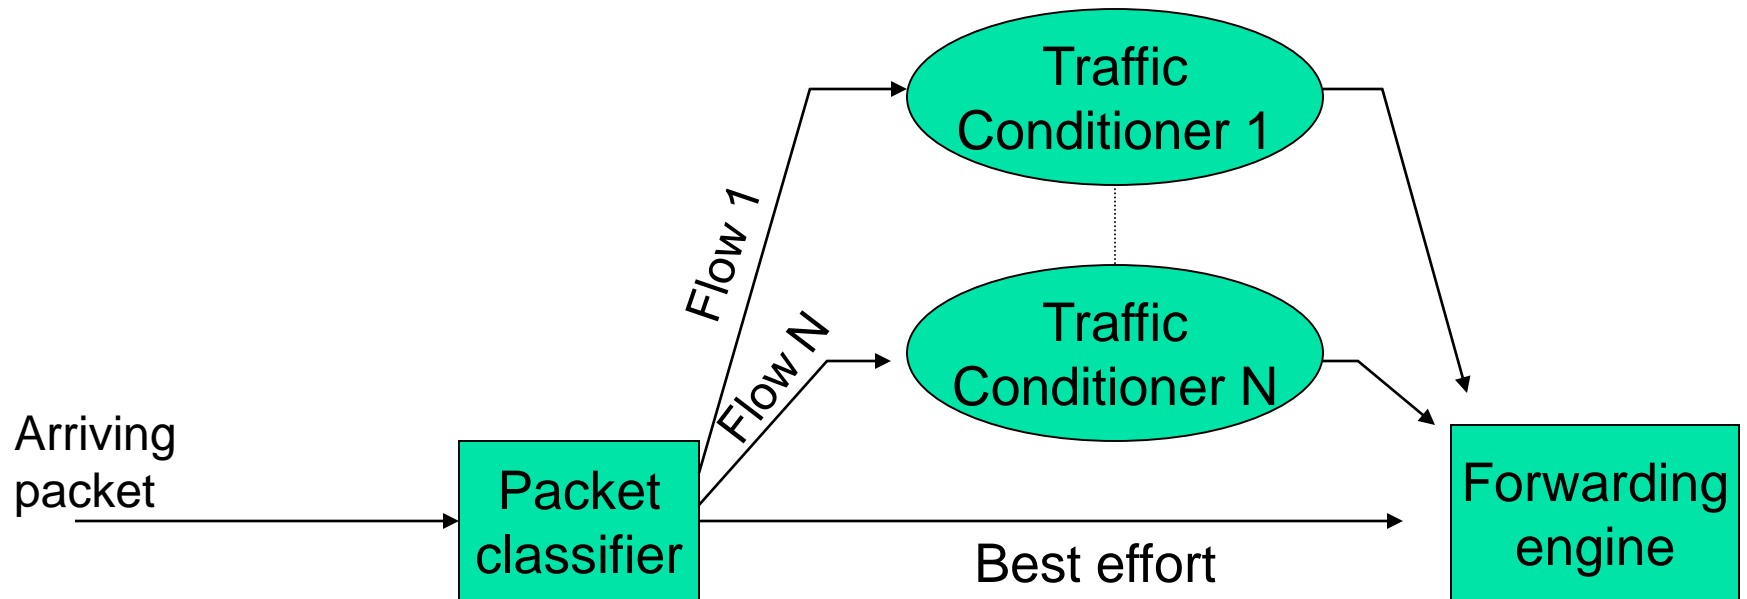
---

- Similar to RED, but with two separate probability curves
- Has two classes, “In” and “Out” (of profile)
- “Out” class has lower  $\text{Min}_{\text{thresh}}$ , so packets are dropped from this class first
  - Based on queue length of all packets
- As avg queue length increases, “in” packets are also dropped
  - Based on queue length of only “in” packets

# RIO Drop Probabilities

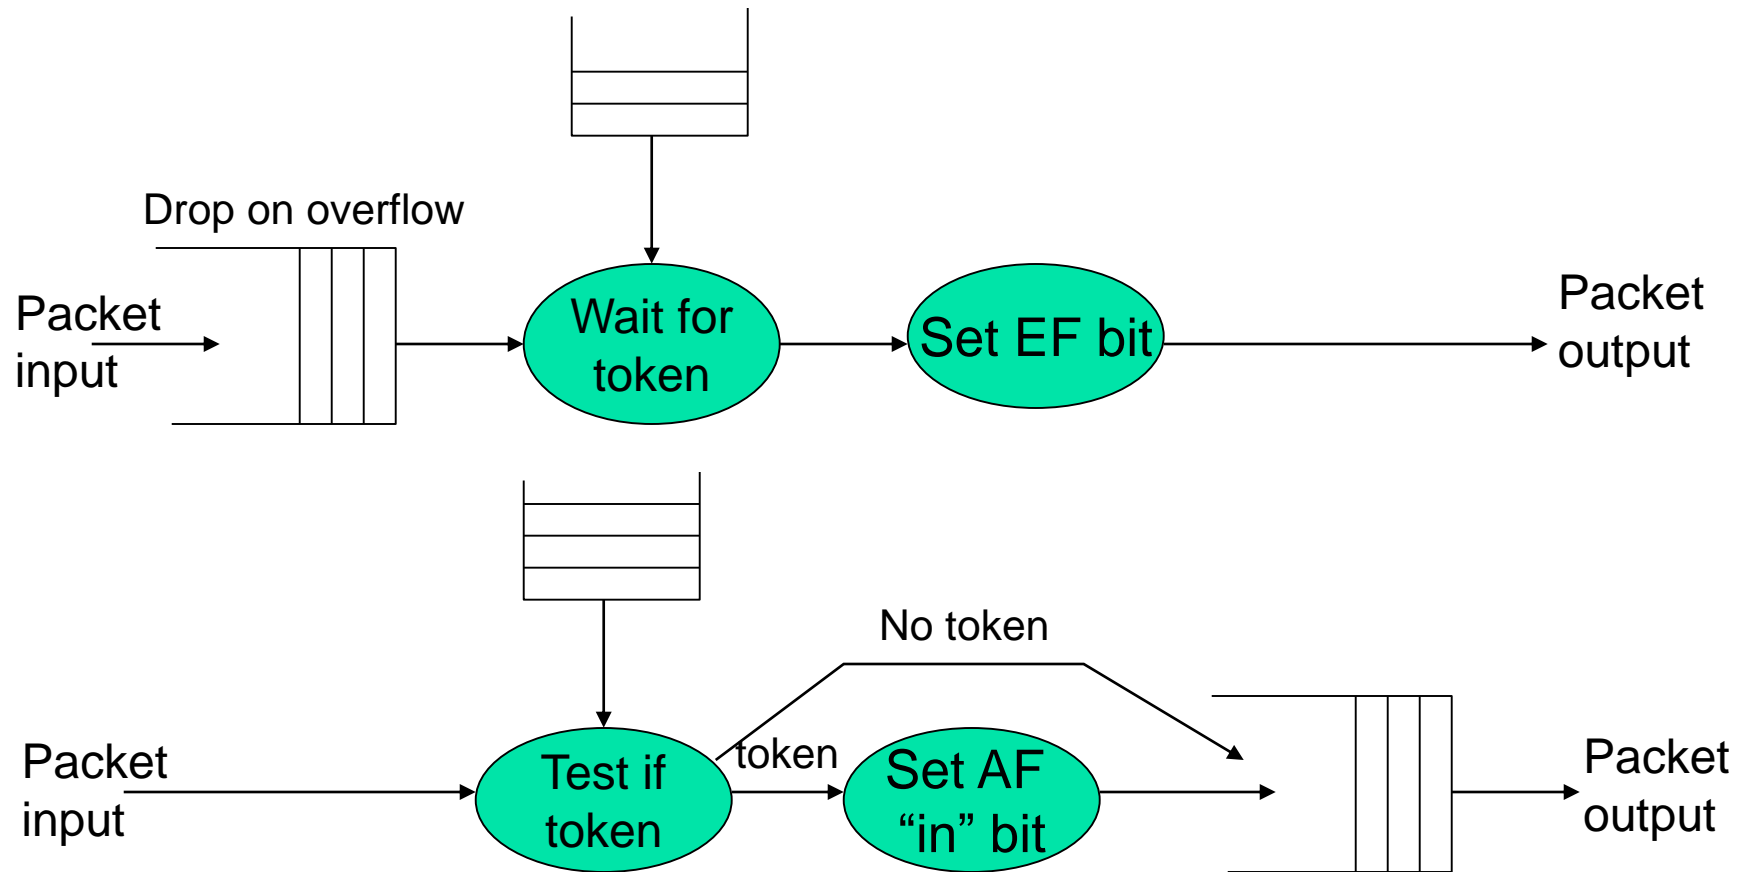


# Edge Router Input Functionality



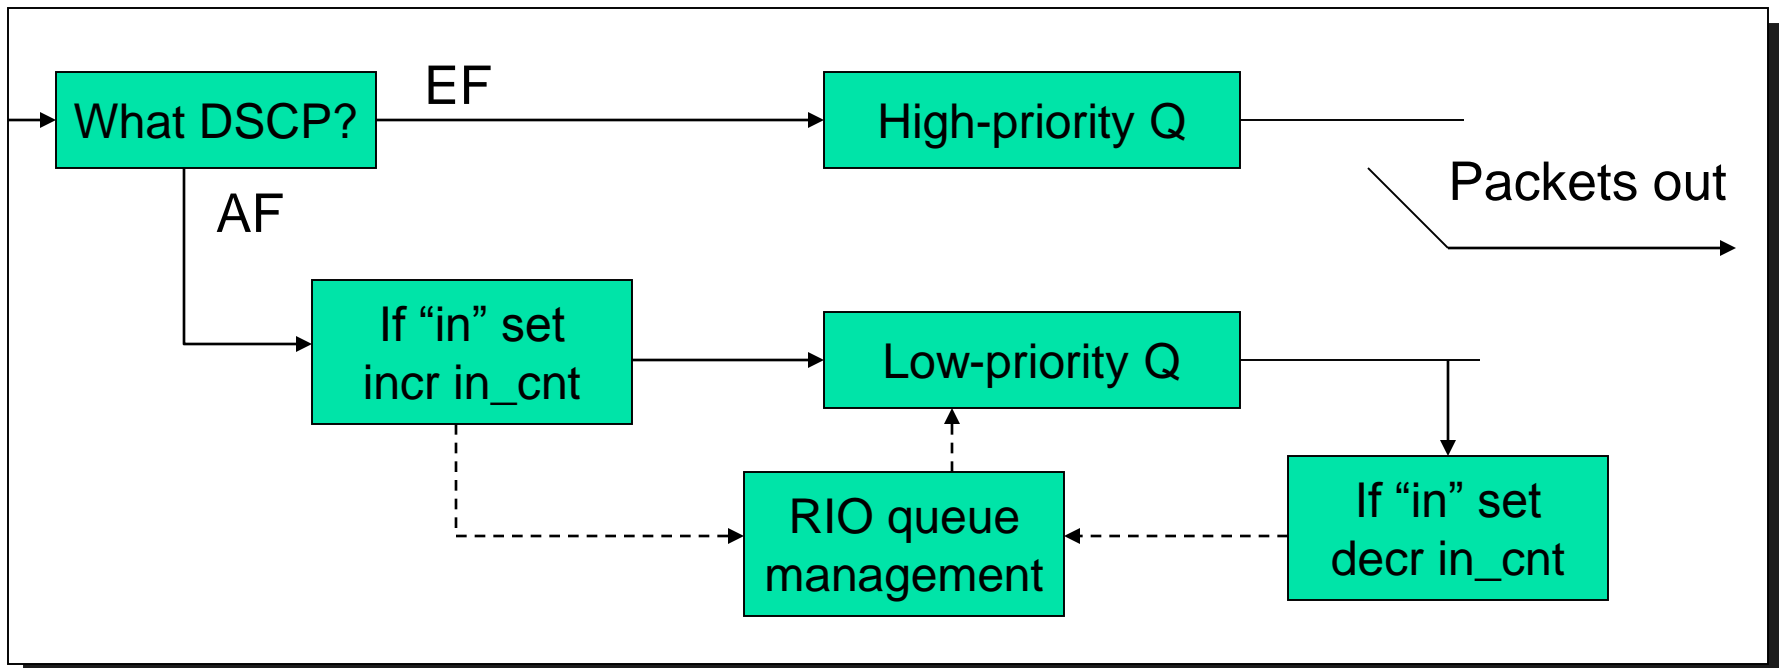
classify packets based on packet header

# Traffic Conditioning

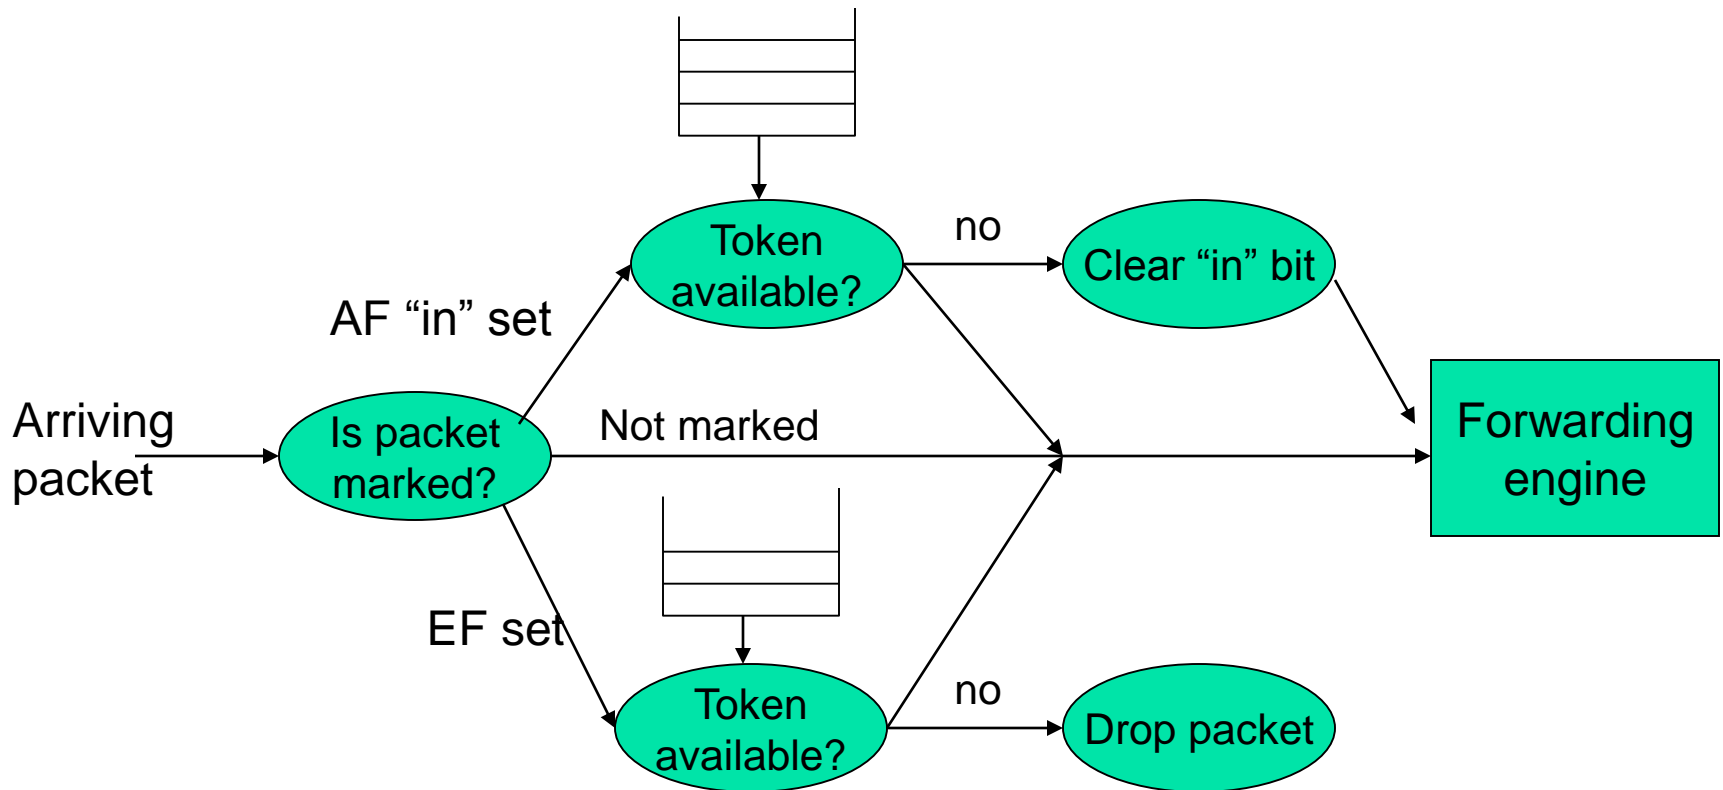


# Router Output Processing

- 2 queues: EF packets on higher priority queue
- Lower priority queue implements RED “In or Out” scheme (RIO)



# Edge Router Policing







# Comparison

	Best-Effort	Diffserv	Intserv
Service	<ul style="list-style-type: none"><li>• Connectivity</li><li>• No isolation</li><li>• No guarantees</li></ul>	<ul style="list-style-type: none"><li>• Per aggregation isolation</li><li>• Per aggregation guarantee</li></ul>	<ul style="list-style-type: none"><li>• Per flow isolation</li><li>• Per flow guarantee</li></ul>
Service Scope	<ul style="list-style-type: none"><li>• End-to-end</li></ul>	<ul style="list-style-type: none"><li>• Domain</li></ul>	<ul style="list-style-type: none"><li>• End-to-end</li></ul>
Complexity	<ul style="list-style-type: none"><li>• No set-up</li></ul>	<ul style="list-style-type: none"><li>• Long term setup</li></ul>	<ul style="list-style-type: none"><li>• Per flow setup</li></ul>
Scalability	<ul style="list-style-type: none"><li>• Highly scalable</li><li>• (nodes maintain only routing state)</li></ul>	<ul style="list-style-type: none"><li>• Scalable (edge routers maintains per aggregate state; core routers per class state)</li></ul>	<ul style="list-style-type: none"><li>• Not scalable (each router maintains per flow state)</li></ul>