



# Advanced Computer Networks

---

Network Virtualization, NFV, SDN



# Outline

---

- Introduction and Historical Perspective
- Reference Model, Architectural Principles and Objectives
- Network Virtualization Projects
- Challenges and Future Directions




# Virtualization As Resource Abstraction

---

- **Virtualization**: transparent **abstraction** of the physical computing platform and resources that supports **multiple** logical views of their properties
- Virtual anything
  - virtual memory/queue
  - OS platform → virtual machines
  - storage → virtual SAN
  - computing → virtual data centers

# Network Virtualization

- Single physical network **appears as** multiple logical networks





# Network Virtualization Environment

---

- Allows the creation of multiple virtual networks on same substrate
- Each virtual network:
  - is a collection of virtual nodes and virtual links
  - is a subset of underlying physical network resources
  - co-exists with, but is isolated from, other virtual networks





# Historical Perspective

---

- Virtual LAN (VLAN)
- Virtual service network (VSN)
- Virtual private network (VPN)
- Active and programmable networks
- Overlay networks


# Virtual LANs

- L2 constructs, share same physical LAN infrastructure
- Groups of hosts with common interests, regardless of connectivity
- Segmented within boundary of broadcast domains → **frame coloring**
- High levels of trust, security, isolation, easy to configure



# Virtual Service Networks


- Generalization of VLAN concept to broader network
- Define multiple network instances that
  - share the physical resources
  - are properly segmented (functionality, access permissions, etc.)





# Virtual Private Networks

- Connect multiple sites using **tunnels** over public network
- L3/L2/L1 technologies
- Intranet vs. extranet
- Limited customization





# Active and Programmable Networks

---

- Support co-existing networks through programmability
- Customized network functionality through
  - programmable interfaces to enable access to switches/routers
  - active code
    - call functions already installed
    - allow user to execute own code at switches/routers
- Operators must open networks to third parties → not used




# Overlay Networks (1)

---

- A network built on top of one or more existing networks
  - without deploying entirely new networking equipments
- Adds an additional layer
- Changes properties in one or more layers of underlying network
- Many networks after PSTN have begun as overlay networks
  - the Internet also starts as an overlay network

# Overlay Networks (2)

- Application layer virtual networks
- Built upon existing network using tunneling/encapsulation
- Implement new services without changes in infrastructure
- Application-specific, not flexible enough





# VN Motivation (1)

---

- The Internet is **ossified**:
  - significant innovation at application and link layers
  - only incremental changes to core protocols
    - half-layer solutions (MPLS, IPSec)
    - adapt to new contexts (TCP over wireless)
    - application overlays (CDN, ALM, OverQoS, SOS, P2P, . . .)



## VN Motivation (2)

---

- **Clean-slate** architecture impossible to deploy
  - multiple stakeholders
    - conflicting goals/policies! difficult to reach agreement
    - almost no economic incentives to introduce changes
  - defining one-size-fits-all architecture a challenge
    - depends on focus and context! 4 distinct FIA projects
    - future needs unknown, difficult to predict
  - must accommodate installed base

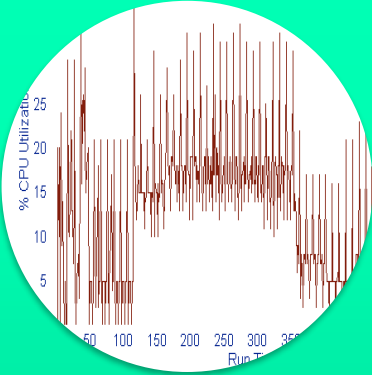


## VN Motivation (3)

---

- Network virtualization offers **all-sizes-fit-into-one** solution
- Open and expandable model
  - multiple heterogeneous architectures on shared physical substrate
  - promotes innovation and customized services/applications
- Testbed for deployment/evaluation of new network architectures and protocols

# Drivers Behind Virtualization



## Hardware Resources Underutilized

- CPU utilizations ~ 10% - 25%
- One server – One Application
- Multi-core even more under-utilized



## Data Centers are running out of space

- Last 10+ years of major server sprawl
- Exponential data growth
- Server consolidation projects just a start



## Rising Energy Costs

- As much as 50% of the IT budget
- In the realm of the CFO and Facilities Mgr. now!



## Administration Costs are Increasing

- Number of operators going up
- Number of Management Applications going up

Operational Flexibility





# Other Significant Virtualization Benefits

---

- Some key benefits:
  - Ability to quickly spawn test and development environments
  - Provides failover capabilities to applications that can't do it natively
  - Maximizes utilization of resources (compute & I/O capacity)
  - Server portability (migrate a server from one host to the other)
- Virtualization is **not** limited to servers and OS
  - Network virtualization
  - Storage virtualization
  - Application virtualization
  - Desktop virtualization

Isolation

Roll-Back

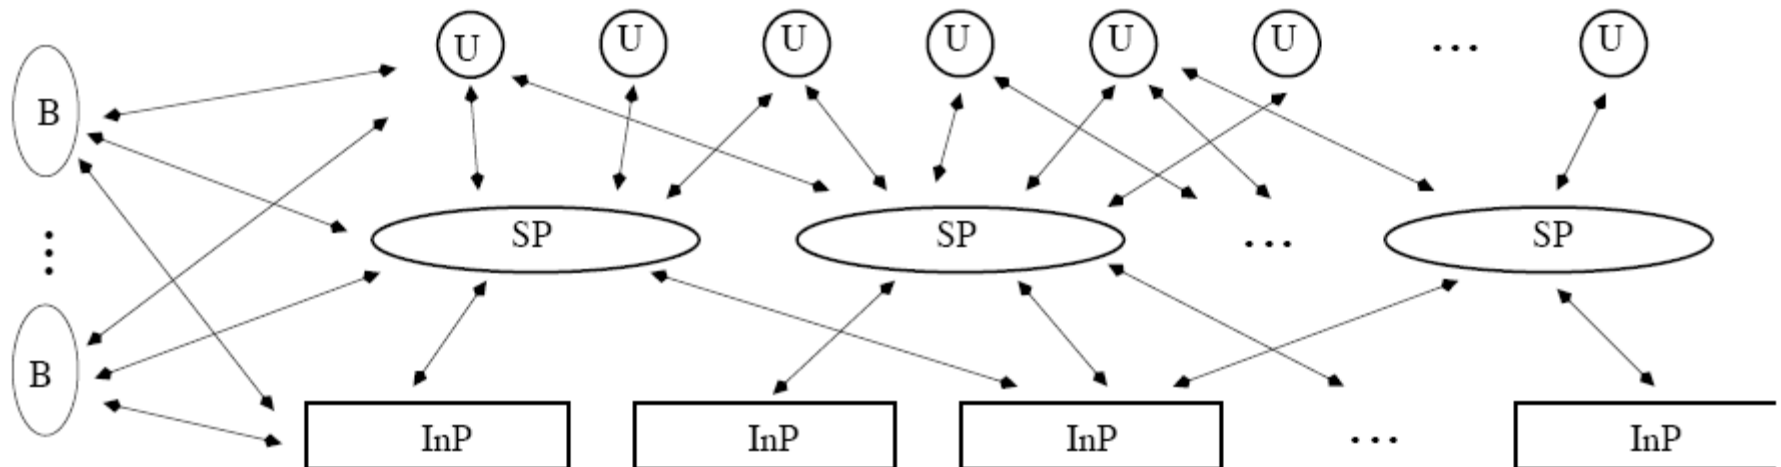
Abstraction

Portability

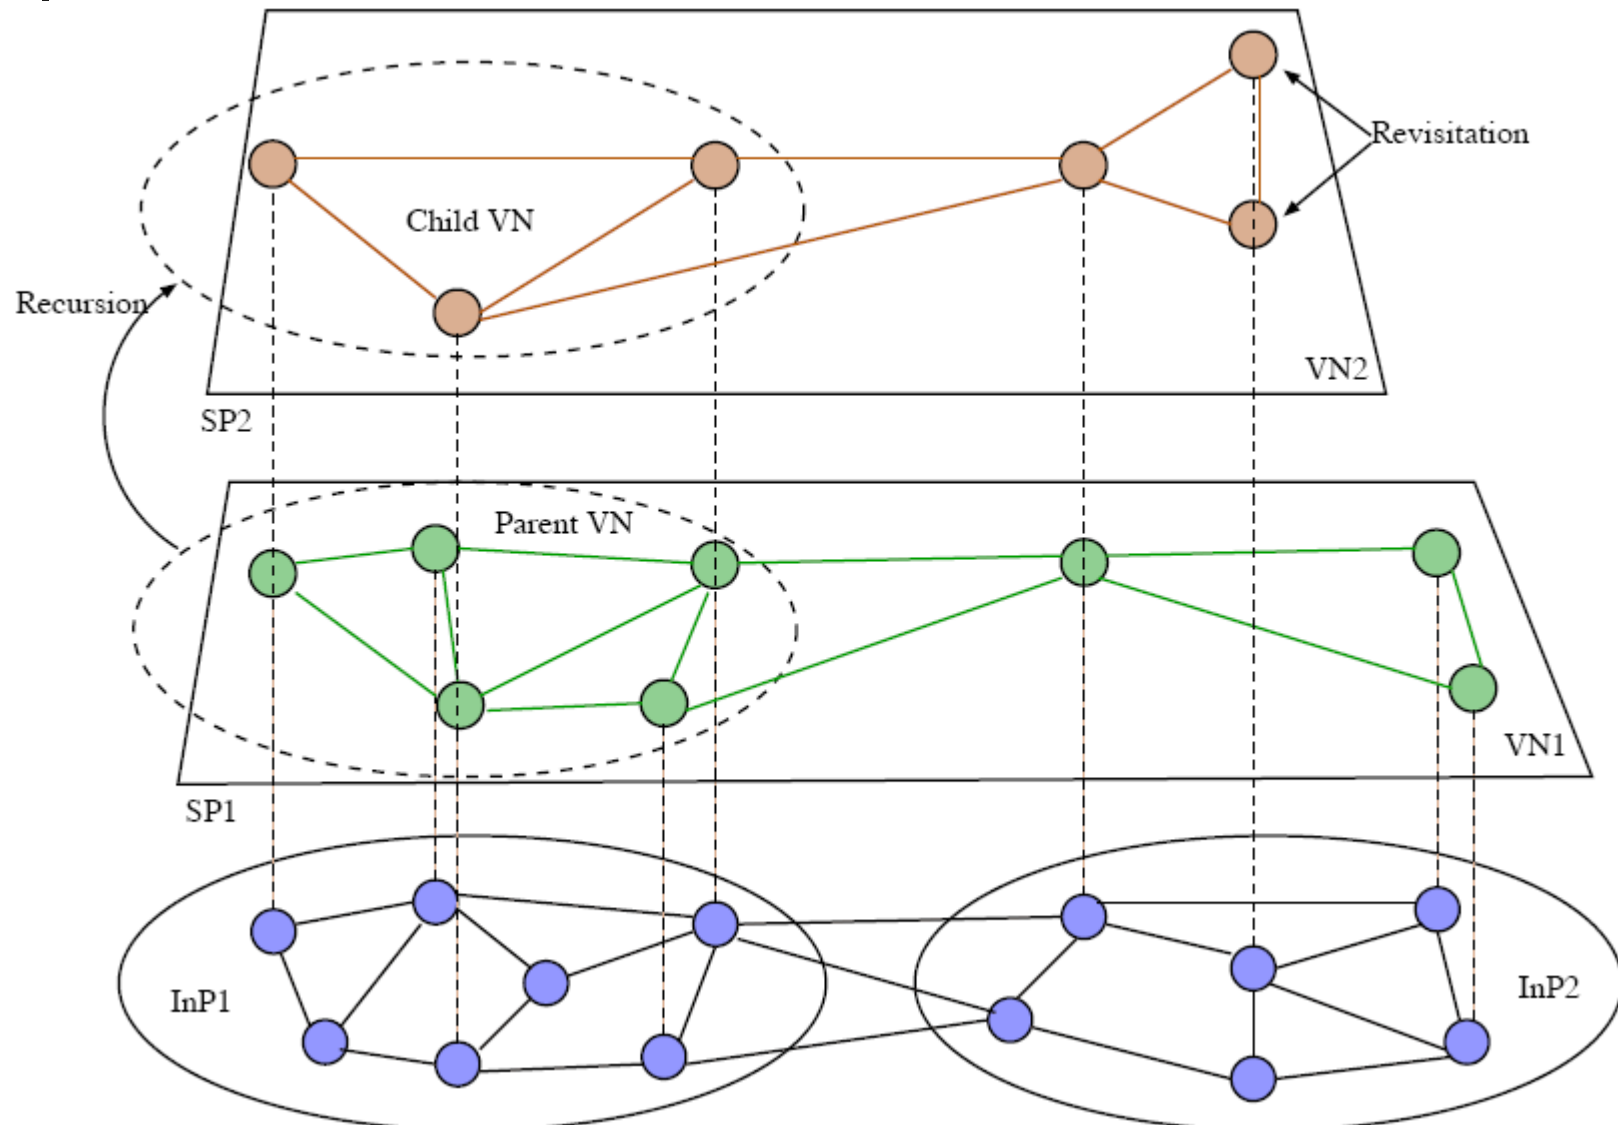
Deployment

# Reference Model

- **Infrastructure providers**: manage physical networks
- **Service providers**:
  - create and manage virtual networks
  - design and deploy end-to-end services
- **Users**: select from competing services offered by providers
- **Brokers**: aggregate offers, match requirements to services/resources



# Architectural Principles (1)





## Architectural Principles (2)

---

- **Concurrency:**
  - multiple service providers compete
  - multiple virtual networks coexist → diversity
- **Nesting:** virtual network hierarchy
  - child VNs derived from parent VNs
- **Inheritance:** child VN inherits architectural attributes
  - creation of value-added services
  - VN economics
- **Revisitation:** single physical node hosts multiple virtual nodes
  - deploy diverse functionalities
  - simplify operation and management



# Design Objectives (1)

---

## ■ Isolation

- in terms of logical view, resources, operation
- attacks, failures, bugs, misconfigurations do not affect other VNs.

## ■ Flexibility

- service providers may assemble VNs with customized
  - network topology
- control and data plane functionality without the need to coordinate with each other

## ■ Scalability

- support multiple VNs
- utilize resources efficiently



## Design Objectives (2)

---

### ■ Programmability

- deploy customized protocols/services in network elements
- at packet level or below (e.g., optical devices?)
- how to expose to service providers/users
- opens up vulnerabilities (active networks redux?)

### ■ Heterogeneity

- support a diverse set of
  - physical network technologies (wireless, sensor, optical, . . . )
  - virtual network capabilities

### ■ Manageability

- Mechanism  $\leftrightarrow$  policy
- “know what happened  $\rightarrow$  SPs and InPs held accountable



## Design Objectives (3)

---

### ■ Dual use

- experimental **and** deployment facility
- design, evaluate, and deploy new services
- plausible pathway for adopting radically new network architectures
- PlanetLab, GENI, VINI

### ■ Legacy support

- existing Internet: another instance of a VN
- deploy IPv6 as VN → compete with IPv4



# Network Virtualization Projects: Classification

---

- Targeted technology
  - wired, wireless, IP, heterogeneous
- Layer (in network stack) where virtualization is introduced
  - Physical → application
- Application domain
  - specific problem domain addressed! wide range
- Virtualization granularity
  - node, link, full





# The Holy Grail of Network Virtualization

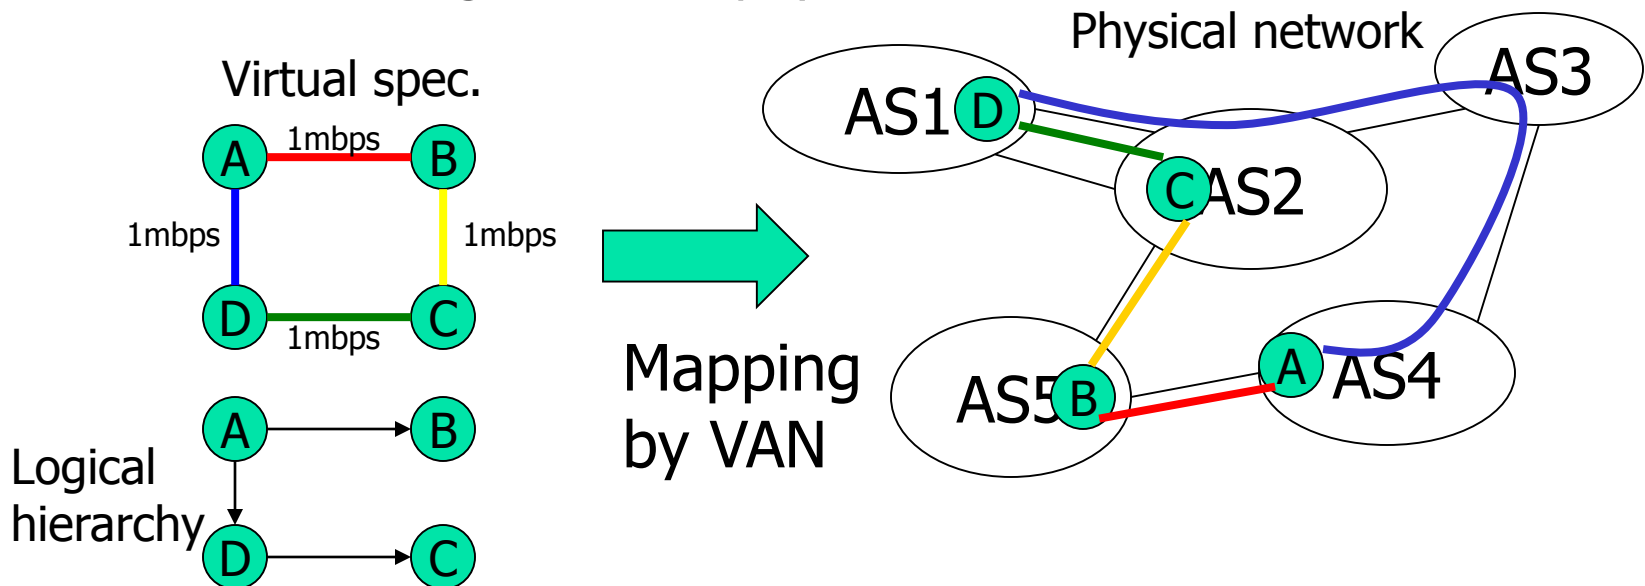
---

- A **network environment** in which multiple SPs:
  - **lease** underlying physical resources from multiple InPs
  - dynamically **compose** heterogeneous VNs
    - co-exist in isolation within same physical infrastructure
  - **compete** with each other by
    - deploying customized E2E services on-the-fly
    - managing the services on the VNs for the end users
    - effectively sharing and utilizing the physical resources

# A Driving Example

A web caching application needs...


- Coverage for certain network area
- Connectivity among caching service components
- Resources to move cached contents
- Solution: requests a VN that provides
  - Coverage: spans a ring between AS1, AS2, AS4, and AS5
  - Resources: provides at least 1mbps for all connections
  - Reliability: prohibits more than 2 virtual links from traversing the same physical link



# Specification Mapping



## Heuristic mapping algorithm

- Sort VNs and PNs by degree; map VN to PN by degree-order
- Mark all PL without enough bandwidth for the VLs as infeasible
- Each PL has a “mapped-onto” counter, initially 0
- pick a VL and map it to a physical path with lowest maximum counter among all PLs traversed
- After each VL is mapped, increment counter and subtract available bandwidth for each PL; mark a PL infeasible as appropriate
- Repeat until all VLs are mapped



# Failure Recovery

- When a physical link fails, the VLs it carries must be restored
  - First try Local repair
    - Find an alternative path with adequate resources between the two disconnected AS'es
    - Fast, and preserve original topology
    - But
      - May violate reliability constraint
      - Alternative path may not exist
  - Example
    - Physical link between 2 and 3 goes down
    - Alternative path goes through 2, 1, 4, and 3




# Virtual Network Embedding

- Virtual Network:

- the logical network of SP


- Substrate Network

- the physical network of InP






# Network Flow Construction

## Virtual Network



## Substrate Network



Auxiliary edge:  
connect a virtual node  
to **eligible** substrate  
nodes

# The NFV Concept

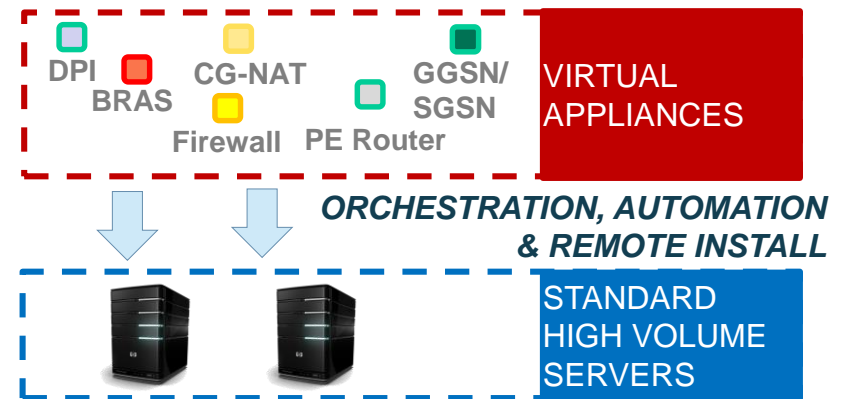
A means to make the **network more flexible and simple by minimising dependence on HW constraints**

## Traditional Network Model: APPLIANCE APPROACH



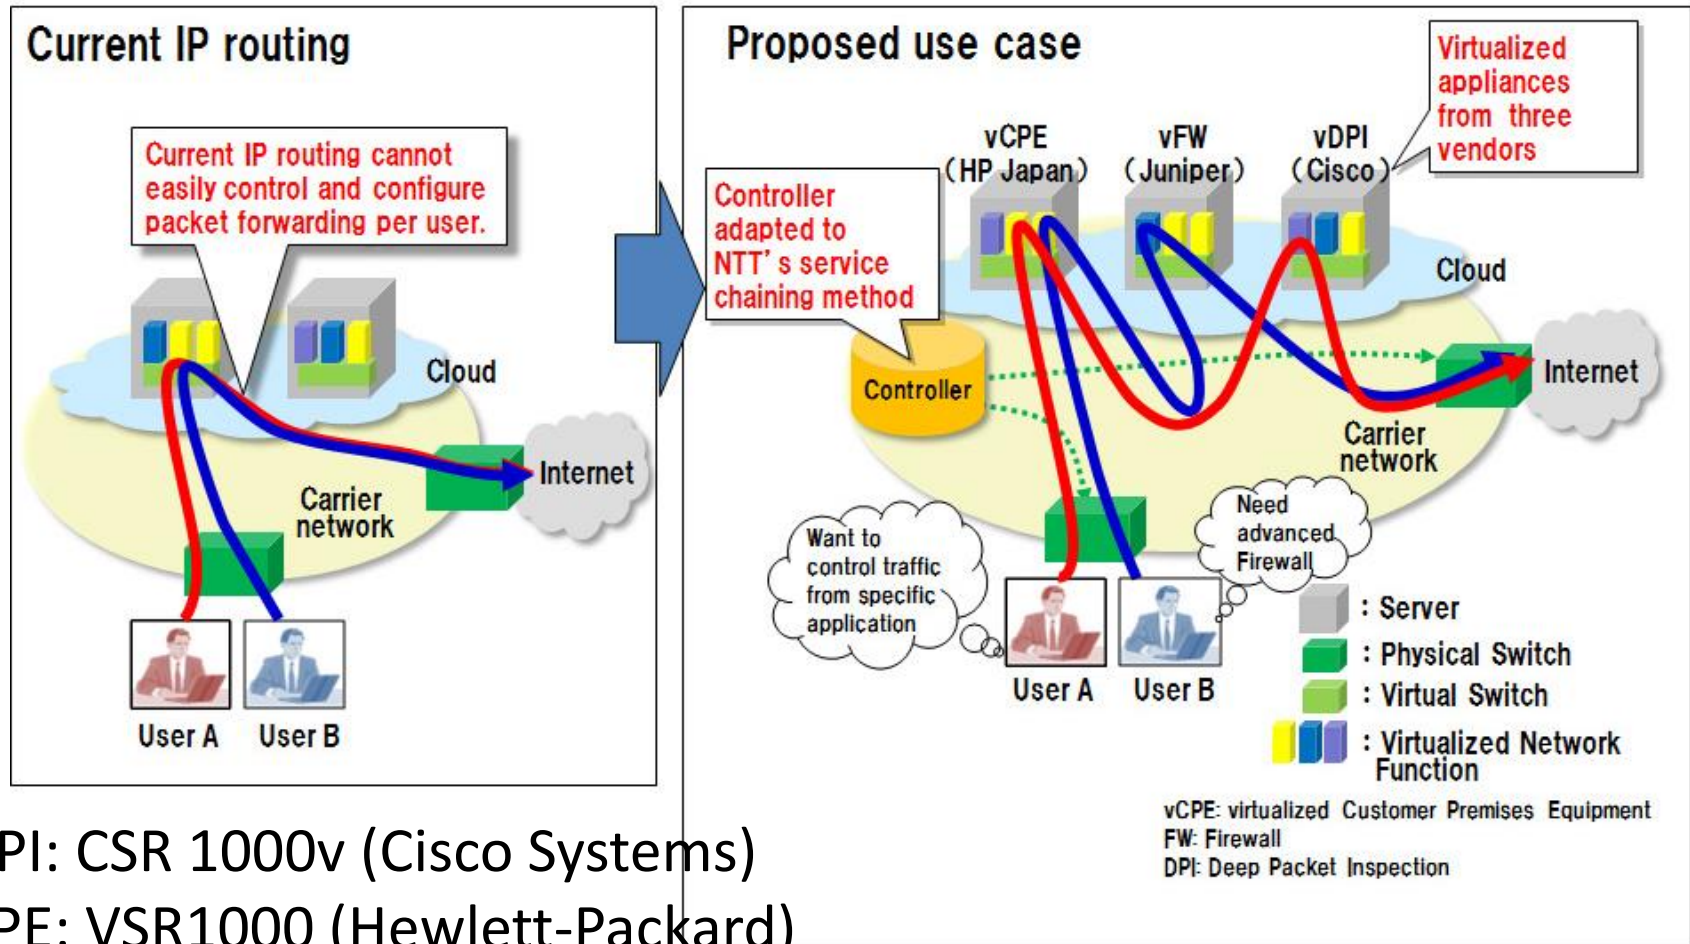
- Network Functions are **based on specific HW&SW**
- **One physical node per role**

## Virtualised Network Model: VIRTUAL APPLIANCE APPROACH



- Network Functions are **SW-based over well-known HW**
- **Multiple roles over same HW**

# Service Chaining for NW Function Selection in Carrier Networks



vDPI: CSR 1000v (Cisco Systems)  
vCPE: VSR1000 (Hewlett-Packard)  
vFW: FireFly (Juniper Networks)  
VIM (NW Controller): Service

Source: ETSI Ongoing






# Rethinking the “Division of Labor”

---

# Traditional Computer Networks


**Control plane:**  
**Distributed algorithms**



**Track topology changes, compute routes, install forwarding rules**

# Traditional Computer Networks

**Management plane:**  
**Human time scale**




Collect measurements and configure  
the equipment



# Death to the Control Plane!

---

- Simpler management
  - No need to “invert” control-plane operations
- Faster pace of innovation
  - Less dependence on vendors and standards
- Easier interoperability
  - Compatibility only in “wire” protocols
- Simpler, cheaper equipment
  - Minimal software




# Software Defined Networking (SDN)

**Logically-centralized control**

**Smart,  
slow**




**API to the data plane  
(e.g., OpenFlow)**



**Dumb,  
fast**

**Switches**

# Controller: Programmability



## Events from switches

Topology changes,  
Traffic statistics,  
Arriving packets

## Commands to switches

(Un)install rules,  
Query statistics,  
Send packets


# E.g.: Network Virtualization

**Controller #1**

**Controller #2**

**Controller #3**

**Partition the space of packet headers**






## A Helpful Analogy

---

From Nick McKeown's talk "Making SDN Work" at the Open Networking Summit, April 2012




# Mainframes




**Vertically integrated**  
**Closed, proprietary**  
**Slow innovation**  
**Small industry**



— Open Interface —



— Open Interface —



**Horizontal**  
**Open interfaces**  
**Rapid innovation**  
**Huge industry**

# Routers/Switches



**Vertically integrated**  
**Closed, proprietary**  
**Slow innovation**



— Open Interface —




— Open Interface —



**Horizontal**  
**Open interfaces**  
**Rapid innovation**

# Challenges: Controller Delay and Overhead

- Controller is much slower than the switch
- Processing packets leads to delay and overhead
- Need to keep most packets in the “fast path”



# Challenges: Distributed Controller

