

CSc 8530

Parallel Algorithms

Spring 2019

January 15th, 2019

Dramatis personae

- **Instructor:** Rolando Estrada
 - Email: restrada1@gsu.edu
 - Office: 1 Park Place, Suite 634
 - Office hours: Tues. 11:30am – 12:30pm or by appointment
- **Grader:** Saeid Motevalialamoti
 - Email: smotevalialamoti1@student.gsu.edu
 - Office: 1 Park Place, Suite 633
 - Office hours: TR 11:00am – 12:00pm or by appointment

Course prerequisites

- CSc 4520/6520 – Design & Analysis of Algorithms
- CSc 4310/6310 – Parallel & Distributed Computing
- Equivalent senior-level course

Conceptual prerequisites

- Standard algorithm analysis
 - Big- O notation, complexity classes, asymptotic bounds, etc.
 - Design techniques, e.g., divide and conquer, greedy algorithms, etc.
- Basic graph theory
 - Graph algorithms: BFS, DFS, Dijkstra, Kruskal, etc.
 - Graph properties: sparse vs. dense, subtypes (regular, small-world, etc.), explicit vs. implicit representations, etc.
- Some programming experience
 - Preferably C/C++
 - Hardware-level a plus

Required textbooks

- *An Introduction to Parallel Algorithms, 1st edition* by Joseph JáJá, Addison-Wesley, 1992
- *Programming Massively Parallel Processors: A Hands-on Approach, 3rd edition* by David B. Kirk and Wen-mei W. Hwu, Morgan Kaufmann, 2016
- Both books are available on Amazon and similar retailers

Course content (general overview)

- Asymptotic analysis of parallel algorithms
- Parallel data handling and message passing
- Parallel algorithm design techniques (e.g., divide and conquer)
- Parallel graph algorithms
- CUDA/GPU programming

Course Grade

- **Rubric**


- Homework 20%
 - Midterm 20%
 - Paper presentation 15%
 - Final project 45%
- The lowest homework grade will be dropped.
 - Grades on group projects will be assigned individually
 - Letter grades will be assigned relative to class performance.

Policies

- Late/missing assignments will not be graded
- Complete academic honesty is expected
 - No cheating
 - No plagiarism
- Cell phones must be off/silent during class
 - No texting, social media, etc.


O -notation

- $O(g(n)) = \{f(n) : \text{there exist positive constants } c \text{ and } n_0 \text{ such that } 0 \leq f(n) \leq cg(n) \text{ for all } n \geq n_0\}$
- $g(n)$ is an *asymptotic upper bound* for $f(n)$
- If $f(n) \in O(g(n))$, we write $f(n) = O(g(n))$
 - **Abuse of notation for convenience**
 - Similarly for the other notations




Ω -notation

- $\Omega(g(n)) = \{f(n) : \text{there exist positive constants } c \text{ and } n_0 \text{ such that } 0 \leq cg(n) \leq f(n) \text{ for all } n \leq n_0\}$
- $g(n)$ is an *asymptotic lower bound* for $f(n)$



Θ -notation

- $\Theta(g(n)) = \{f(n) : \text{there exist positive constants } c_1, c_2, \text{ and } n_0 \text{ such that } 0 \leq c_1g(n) \leq f(n) \leq c_2g(n) \text{ for all } n \leq n_0\}$
- $g(n)$ is an *asymptotically tight bound* for $f(n)$
- **Theorem:** $f(n) = \Theta(g(n))$ if and only if $f = O(g(n))$ and $f = \Omega(g(n))$



Asymptotic notation in equations

- **On right-hand side:** $O(g(n))$ stands for some *anonymous function* in the set $O(g(n))$
- E.g.: $2n^2 + 3n + 1 = 2n^2 + \Theta(n)$

Asymptotic notation in equations

- **On right-hand side:** $O(g(n))$ stands for some *anonymous function* in the set $O(g(n))$
- E.g.: $2n^2 + 3n + 1 = 2n^2 + \Theta(n)$
 - Means $2n^2 + f(n)$, with $f(n) \in \Theta(n)$

Asymptotic notation in equations

- **On right-hand side:** $O(g(n))$ stands for some *anonymous function* in the set $O(g(n))$
- E.g.: $2n^2 + 3n + 1 = 2n^2 + \Theta(n)$
 - Means $2n^2 + f(n)$, with $f(n) \in \Theta(n)$
 - Here, $f(n) = 3n + 1$

Asymptotic notation in equations

- **On right-hand side:** $O(g(n))$ stands for some *anonymous function* in the set $O(g(n))$
- E.g.: $2n^2 + 3n + 1 = 2n^2 + \Theta(n)$
 - Means $2n^2 + f(n)$, with $f(n) \in \Theta(n)$
 - Here, $f(n) = 3n + 1$
 - But, $f(n) = 4n + 5$ or $f(n) = 10000n + 1000$, etc. are all equivalent from an asymptotic perspective

Asymptotic notation in equations

- **On right-hand side:** $O(g(n))$ stands for some *anonymous function* in the set $O(g(n))$
- E.g.: $2n^2 + 3n + 1 = 2n^2 + \Theta(n)$
 - Means $2n^2 + f(n)$, with $f(n) \in \Theta(n)$
 - Here, $f(n) = 3n + 1$
 - But, $f(n) = 4n + 5$ or $f(n) = 10000n + 1000$, etc. are all equivalent from an asymptotic perspective
- **On left-hand side:** there is a way to choose the corresponding anonymous functions on the *right-hand side* to make the equation valid
- Interpret $2n^2 + \Theta(n) = \Theta(n^2)$

Asymptotic notation in equations

- **On right-hand side:** $O(g(n))$ stands for some *anonymous function* in the set $O(g(n))$
- E.g.: $2n^2 + 3n + 1 = 2n^2 + \Theta(n)$
 - Means $2n^2 + f(n)$, with $f(n) \in \Theta(n)$
 - Here, $f(n) = 3n + 1$
 - But, $f(n) = 4n + 5$ or $f(n) = 10000n + 1000$, etc. are all equivalent from an asymptotic perspective
- **On left-hand side:** there is a way to choose the corresponding anonymous functions on the *right-hand side* to make the equation valid
- Interpret $2n^2 + \Theta(n) = \Theta(n^2)$
 - For all functions $f(n) \in \Theta(n)$, there exists a function $g(n) \in \Theta(n^2)$, such that $2n^2 + f(n) = g(n)$

Important classes of algorithms

- Most algorithms you'll encounter in practice are either:
 - **Constant time:** $\Theta(1)$
 - Running time is independent of input size
 - **Logarithmic time:** $\Theta(\log(n))$
 - Running time is proportional to the *number of bits* needed to encode the input
 - **Linear time:** $\Theta(n)$
 - Running time is proportional to the input size
 - **Log-linear time:** $\Theta(n \log(n))$
 - How many times we execute an $\Theta(n)$ operation depends on the input size's number of bits
 - **Polynomial time:** $\Theta(n^p)$
 - Common cases: $\Theta(n)$, $\Theta(n^2)$ (quadratic), $\Theta(n^3)$ (cubic)
 - Running time is proportional to a number of subsets (e.g., pairs for quadratic, triples for cubic, etc.)
 - **Exponential time:** $\Theta(2^{n^p})$
 - Common case: $\Theta(2^n)$
 - Running time *doubles* every time the input size grows by one
 - Practical only for small inputs

Graph definition

- A graph $G = (V, E)$ is defined by two sets:
 - A set of n vertices V (also called nodes)
 - A set of m edges E (also called links)
- All the elements in both V and E are unique (i.e., no repeated values)
- Every edge $e = (u, v) \in E$ is a tuple (i.e., two *ordered* values), such that $u, v \in V$
 - In other words, each edge is defined by its starting and ending vertices
- In general, $m = O(n^2)$ (why?)

Graph terminology

- A graph is **undirected**, if and only if

$$e = (u, v) \in E \Rightarrow (v, u) \in E$$

- In other words, in an undirected graph, if it is possible to go from u to v , then it is also possible to go from v to u .
- A graph is **directed** if the above condition does **not** hold
- The vertex v is a **neighbor** of the vertex u if and only if $(u, v) \in E$.
 - For an undirected graph: if v is a neighbor of u , then u is also a neighbor of v .
- A **self-loop** is an edge where the starting and ending vertices are the same, i.e., $e = (v, v)$
- A **path** is a sequence of vertices (v_1, v_2, \dots, v_k) , such that v_{i+1} is a neighbor of v_i .

Graph representations


- There are two standard ways of representing graphs in a computer. For both, we assume that the vertices are indexed $1, \dots, n$ in an arbitrary manner
 - **Adjacency list:** Each vertex $v \in V$ has an associated array Adj_v that lists all of its neighbors (randomly or ordered by index)
 - **Adjacency matrix:** We define an *adjacency matrix* A , such that

$$a_{i,j} = \begin{cases} 1 & \text{if } (i,j) \in E \\ 0 & \text{otherwise.} \end{cases}$$


Graph representations

- Adjacency lists are usually better for **sparse** graphs (graphs where the number of edges is similar to the number of vertices, $m = \Theta(n)$)
- Adjacency matrices are sometimes more convenient for **dense** graphs, where $m = \Theta(n^2)$.
- More generally, though, some operations are easier on one representation than another. For example:
 - Determining if v is a neighbor of u takes $\Theta(1)$ for adjacency matrices, but $\Theta(n)$ for adjacency lists
 - Conversely, finding the number of neighbors of u takes $\Theta(1)$ for adjacency lists, but $\Theta(n)$ for adjacency matrices
- **The best representation will depend on your problem**

Undirected graph



(a)




(b)


	1	2	3	4	5
1	0	1	0	0	1
2	1	0	1	1	1
3	0	1	0	1	0
4	0	1	1	0	1
5	1	1	0	1	0

(c)

Directed graph



(a)



(b)

	1	2	3	4	5	6
1	0	1	0	1	0	0
2	0	0	0	0	1	0
3	0	0	0	0	1	1
4	0	1	0	0	0	0
5	0	0	0	1	0	0
6	0	0	0	0	0	1

(c)