

# CSc 8530 Parallel Algorithms

Spring 2019

January 22nd, 2019

# Divide and conquer

- There are many strategies for designing algorithms
- The most basic is **incremental**
  - We iteratively add one element to our partial solution at a time until we have a complete solution
  - Example: insertion sort
- **Divide and conquer** is another common approach
  - **Divide** the problem into smaller instances of the same problem
  - **Conquer** the subproblems by solving them *recursively*
    - **Base case:** if a subproblem is small enough, solve it directly
  - **Combine** the solutions to the subproblems to solve the original problem
  - Example: merge sort

# Greedy algorithms

- Algorithms for optimizing a value (e.g., the minimum cost of a spanning tree) typically go through a sequence of steps, with a set of choices at each step
- An algorithm is **greedy** if, when faced with a set of possible actions, it always picks the one that looks best at the moment
  - It doesn't factor in how earlier choices influence later ones
- We make a *locally optimal choice* in the hope of getting a *globally optimal solution*

# Dynamic programming

- **Dynamic programming (DP)** is a powerful optimization technique which breaks a problem into subproblems
  - Similar to divide-and-conquer, but DP caches intermediate results
    - Avoids solving the same subproblem twice
  - Similar to greedy algorithms, but applies to problems where we have to factor in the subsequent cost of an action
    - In the greedy case, we only care about the local, immediate cost
- **Note:** the term “programming” refers to scheduling, not code
- As in the phrases: “Today’s reception has been programmed for 5:00pm” or “Get with the program”

# Pointer arithmetic

- Pointers can be manipulated like other integers
  - Addition, subtraction, multiplication, etc.
- Pointers often point to the start of a data structure (e.g., an array)
- We usually add a constant to a pointer to access different parts of that structure
- Example:  $p = p + 1$  accesses the adjacent memory location
- **A point of caution:** Trying to access memory locations outside of your program's valid area will result in a segmentation fault

```
// Variable of type int
int x = 5;

// pnt points to x. & is called a
// derefencing (or address) operator
int *pnt = &x;

// y has the same value as x (5)
// The * operator access the value
// that the pointer points to
int y = *pnt;

// Pointer z points to x
int *z = pnt;

// How are x, y, and z affected?
*pnt = 10;
```

# Dynamic memory allocation

```
// Static memory allocation
int array[10];

// Dynamic memory allocation
// with error checking
int *array = malloc(size*sizeof(int));
if (array == NULL) {
    fprintf(stderr, "malloc failed\n");
    return(-1);
}

// ... use the array

// Deallocate the memory
free(array);
```

# Parallel speedup

- Let  $P$  be a computational problem with inputs of size  $n$
- We denote the best-possible sequential (i.e., classic) complexity of  $P$  as  $T^*(n)$
- Let  $A$  be a parallel algorithm that solves  $P$  in time  $T_p(n)$  using  $p$  processors
- Then, the **speedup** achieved by  $A$  is:

$$S_p(n) = \frac{T^*(n)}{T_p(n)}$$

- By construction,  $S_p(n) \leq p$
- We would like  $S_p(n) \approx p$ 
  - i.e., each processor should do around  $1/p$  of the work of a single one
- In practice, inefficiencies in concurrency, synchronization, communication, etc. reduce the actual speedup

# Parallel efficiency

- The **efficiency** of a parallel algorithm  $A$  is given by:

$$E_p(n) = \frac{T_1(n)}{pT_p(n)}$$

- $T_1(n)$  is the running time of the parallel algorithm with a single processor
  - Not necessarily equal to  $T^*(n)$
- Efficiency measures how much bang for our buck we get per processor
- Ideally,  $E_p(n) \approx 1$
- Again, inefficiencies reduce this value in practice



# Computational models

- The **random-access memory** (RAM) model is the standard for sequential algorithms
  - Assumes a single processor and that memory can be accessed in constant time
- No parallel model has the same level of consensus
- Such a model faces two conflicting requirements:
  - **Simplicity:** easy to analyze and hardware-independent
  - **Implementability:** pseudocode should easily translate to implementable algorithms
- Here, we will introduce three simple parallel models:
  - Directed acyclic graphs (dags)
  - The shared-memory model
  - The network model
- In the bulk of this course, we will use the shared-memory model for analysis


# Directed acyclic graphs

- Directed acyclic graphs (dags) are a special case of directed graphs
  - Have **no** directed cycles
    - i.e., it is not possible to return to a node (go backwards)
  - Suitable for modeling **causal processes** and **precedence relations**
- **Some terminology:**
  - *in-degree*: number of incoming edges to a node
  - *out-degree*: number of outgoing edges
  - node  $u$  is the *parent* (*ancestor*) of  $v$  iff there is an edge (path) from  $u$  to  $v$
  - Conversely,  $v$  is called the child (descendant) of  $u$

# Dags for parallel processing

- We can represent computations using dags
  - Nodes with zero in-degree are inputs (also called **leafs**)
  - Nodes with zero out-degree are outputs (also called **roots** or **sinks**)
  - For simplicity, here we assume all internal vertices have in-degree  $\leq 2$
- Each node represents an  $O(1)$  (constant-time) operation
- This model is best-suited for numerical computations
- For simplicity, we will assume no loops (what the book strangely calls branching)
  - We can always unroll a loop by duplicating it the appropriate number of times
- Node order represents **precedence**
  - What operations must come before and after

# Parallel sums



- Two dags for computing the sum  $S$  of the  $n = 2^k$  elements of an array  $A$
- Note how the depths (max distance from leafs to root) differ significantly between the two choices:  $O(n)$  vs  $O(\log(n))$

# Dag model

- In the dag model, we assume that:
  - *Every processor can access the data computed by any other processor without incurring additional cost*
- A particular implementation is defined by **scheduling** each node for execution on a processor
- Given  $p$  processors, we associate a pair  $(j_i, t_i)$  with each internal node  $i$ :
  - $j_i$  is the processor used for node  $i$
  - $t_i$  is the time at which we process node  $i$
- The following two conditions must hold:
  - If  $t_i = t_k$ , for some  $i \neq k$ , then  $j_i \neq j_k$ 
    - Each processor can only process one node at a time
  - If  $(i, k)$  is an edge, then  $t_k \geq t_i + 1$ 
    - Node  $i$  has to be processed before node  $k$


# Dag model

- Input nodes have  $t_i = 0$  and no processor is allocated to them
- The sequence  $\{(j_i, t_i) \mid i \in N\}$  is an execution **schedule**
  - With  $p$  processors
  - $N$  is the number of nodes in the dag
- The time to execute a particular schedule is  $\max_{i \in N} t_i$
- The **parallel complexity** is

$$T_p(n) = \min \{ \max_{i \in N} t_i \}$$


- The minimum is taken over all possible schedules with  $p$  processors
- The depth of the dag is a lower bound on  $T_p(n)$ , for any  $p$

# Parallel sums revisited



- The best schedule for the leftmost algorithm is  $O(n)$ , regardless of  $p$
- The best schedule for the rightmost one is  $O(\log(n))$  with  $n/2$  processors

# Another example: matrix multiplication




- Let  $A$  and  $B$  be two  $n \times n$  matrices
- The standard algorithm for  $C = AB$  computes the terms  $C(i, j) = \sum_{l=1}^n A(i, l)B(l, j)$
- A parallel algorithm with  $n^3$  processors can compute  $C$  in  $O(\log(n))$



# The shared-memory model

- A natural extension of the sequential RAM model
- Many processors have access to a single, **shared memory** unit (also called **global memory**)
- Each processor also has its own **local memory**
- Processors communicate by exchanging data through the shared memory
- Each processor is indexed by a unique id



# The shared-memory model

- In **synchronous** mode, all processors operate in lock-step, under a common clock
- In **asynchronous** mode, each has a independent clock
- Synchronous mode is called the **parallel random-access machine (PRAM)** model
  - It it the model we will primarily study in this class
- Asynchronous mode requires additional checks to make sure the data is up-to-date when accessed
- Both models are **multiple instruction multiple data (MIMD)**
- The amount of **communication** is given by the size of data transferred via the shared memory
- A **global read** ( $X, Y$ ) moves the variable  $X$  into the local memory  $Y$
- A **global write** ( $U, V$ ) does the opposite

## Example: matrix-vector multiplication

- Let  $A$  be an  $n \times n$  matrix and let  $x$  be an  $n$ -dimensional vector
- For simplicity, assume we have  $p \leq n$  processors, such that  $r = n/p$  is an integer
- We also assume *asynchronous* operation

- We partition  $A = \begin{bmatrix} A_1 \\ A_2 \\ \vdots \\ A_p \end{bmatrix}$

- Each block  $A_i$  is of size  $r \times n$
- We can compute  $y = Ax$  as follows:
  - Each processor  $P_i$  reads  $A_i$  and  $x$
  - Computes  $z = A_i x$
  - Writes  $z$  ( $r$  values) to the appropriate location in the shared memory

# Example: matrix-vector multiplication – pseudocode

## ALGORITHM 1.1

### (Matrix Vector Multiplication on the Shared-Memory Model)

**Input:** An  $n \times n$  matrix  $A$  and a vector  $x$  of order  $n$  residing in the shared memory. The initialized local variables are (1) the order  $n$ , (2) the processor number  $i$ , and (3) the number  $p \leq n$  of processors such that  $r = n/p$  is an integer.

**Output:** The components  $(i - 1)r + 1, \dots, ir$  of the vector  $y = Ax$  stored in the shared variable  $y$ .

**begin**

1. **global read**( $x, z$ )
2. **global read**( $A((i - 1)r + 1:ir, 1 : n), B$ )
3. Compute  $w = Bz$ .
4. **global write**( $w, y((i - 1)r + 1 : ir)$ )

**end**

- The above algorithm is specified for a single processor
  - We execute it in parallel on each available one
- The processors can operate asynchronously, because they do not write to overlapping variables or memory locations

# Example: matrix-vector multiplication – pseudocode

## ALGORITHM 1.1

### (Matrix Vector Multiplication on the Shared-Memory Model)

**Input:** An  $n \times n$  matrix  $A$  and a vector  $x$  of order  $n$  residing in the shared memory. The initialized local variables are (1) the order  $n$ , (2) the processor number  $i$ , and (3) the number  $p \leq n$  of processors such that  $r = n/p$  is an integer.

**Output:** The components  $(i - 1)r + 1, \dots, ir$  of the vector  $y = Ax$  stored in the shared variable  $y$ .

**begin**

1. **global read**( $x, z$ )
2. **global read**( $A((i - 1)r + 1 : ir, 1 : n), B$ )
3. Compute  $w = Bz$ .
4. **global write**( $w, y((i - 1)r + 1 : ir)$ )

**end**

- Analysis:
  - Steps 1 and 2 transfer  $O(n^2/p)$  values from the shared memory into each processor
  - Step 3 requires  $O(n^2/p)$  arithmetic operations
  - Step 4 stores  $n/p$  numbers from local to shared memory