

CSc 8530 Parallel Algorithms


Spring 2019

January 24th, 2019

Dags for parallel processing

- We can represent computations using dags
 - Nodes with zero in-degree are inputs (also called **leafs**)
 - Nodes with zero out-degree are outputs (also called **roots** or **sinks**)
 - For simplicity, here we assume all internal vertices have in-degree ≤ 2
- Each node represents an $O(1)$ (constant-time) operation
- This model is best-suited for numerical computations
- For simplicity, we will assume no loops (what the book strangely calls branching)
 - We can always unroll a loop by duplicating it the appropriate number of times
- Node order represents **precedence**
 - What operations must come before and after

Parallel sums



- Two dags for computing the sum S of the $n = 2^k$ elements of an array A
- Note how the depths (max distance from leafs to root) differ significantly between the two choices: $O(n)$ vs $O(\log(n))$

Dag model


- Input nodes have $t_i = 0$ and no processor is allocated to them
- The sequence $\{(j_i, t_i) \mid i \in N\}$ is an execution **schedule**
 - With p processors
 - N is the number of nodes in the dag
- The time to execute a particular schedule is $\max_{i \in N} t_i$
- The **parallel complexity** is

$$T_p(n) = \min \{ \max_{i \in N} t_i \}$$

- The minimum is taken over all possible schedules with p processors
- The depth of the dag is a lower bound on $T_p(n)$, for any p

The shared-memory model

- A natural extension of the sequential RAM model
- Many processors have access to a single, **shared memory** unit (also called **global memory**)
- Each processor also has its own **local memory**
- Processors communicate by exchanging data through the shared memory
- Each processor is indexed by a unique id



Example: matrix-vector multiplication – pseudocode

ALGORITHM 1.1

(Matrix Vector Multiplication on the Shared-Memory Model)

Input: An $n \times n$ matrix A and a vector x of order n residing in the shared memory. The initialized local variables are (1) the order n , (2) the processor number i , and (3) the number $p \leq n$ of processors such that $r = n/p$ is an integer.

Output: The components $(i - 1)r + 1, \dots, ir$ of the vector $y = Ax$ stored in the shared variable y .

begin

1. **global read**(x, z)
2. **global read**($A((i - 1)r + 1 : ir, 1 : n), B$)
3. Compute $w = Bz$.
4. **global write**($w, y((i - 1)r + 1 : ir)$)

end

- Analysis:
 - Steps 1 and 2 transfer $O(n^2/p)$ values from the shared memory into each processor
 - Step 3 requires $O(n^2/p)$ arithmetic operations
 - Step 4 stores n/p numbers from local to shared memory

Example: sum on the PRAM model

- Given an array A with $n = 2^k$ values and a PRAM with p processors
- We wish to compute $S = A[1] + A[2] + \dots A[n]$
- A parallel implementation will run fastest with n processors
- However, not all processors are needed at every iteration

ALGORITHM 1.2

(Sum on the PRAM Model)

Input: An array A of order $n = 2^k$ stored in the shared memory of a PRAM with n processors. The initialized local variables are n and the processor number i .

Output: The sum of the entries of A stored in the shared location S . The array A holds its initial value.

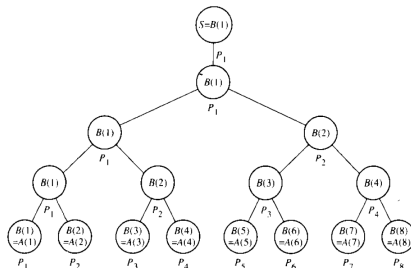
begin

- global read**($A(i), a$)
- global write**($a, B(i)$)
- for** $h = 1$ **to** $\log n$ **do**
 if ($i \leq n/2^h$) **then**
 begin
 global read($B(2i - 1), x$)
 global read($B(2i), y$)
 Set $z := x + y$
 global write($z, B(i)$)
 end
- if** $i = 1$ **then** **global write**(z, S)

end

Example: sum on the PRAM model

- Given an array A with $n = 2^k$ values and a PRAM with p processors
- We wish to compute $S = A[1] + A[2] + \dots A[n]$
- A parallel implementation will run fastest with n processors
- However, not all processors are needed at every iteration



PRAM variations

- PRAM variants differ in how they handle simultaneous access to the same location in shared memory
 - **Exclusive read exclusive write (EREW)**
 - **Concurrent read exclusive write (CREW)**
 - **Concurrent read concurrent write (CRCW)**
- Furthermore, we have three subtypes of CRCW:
 - **Common** CRCW PRAM
 - Allows concurrent writes only when all processors attempt to write *the same value*
 - **Arbitrary** CRCW PRAM
 - Allows an arbitrary processor to succeed
 - **Priority** CRCW PRAM
 - Assumes processors have a priority (based on their ids)
 - The lowest id wins
- EREW, CREW, and, CRCW differ slightly in their computational power
 - i.e., in the space of functions they can theoretically compute

PRAM simplifications

- In the book and our slides, we will omit details concerning memory-access operations
- e.g., an instruction such as $A = B + C$ will really mean:
 - 1 **global read**(B, x)
 - 2 **global read**(C, y)
 - 3 $z = x + y$
 - 4 **global write**(z, A)

The network model


- A **network** is a graph $G = (V, E)$
 - The nodes V are the processors
 - The edges E are two-way communication links between processors
- There is **no** shared memory
 - Each processor does have local memory
- The model can be either **synchronous** or **asynchronous**
- **send**(X, i) instruction: sends X to processor P_i (and continue executing the next instruction immediately)
- **receive**(Y, j) operation: wait for Y from processor P_j (and suspend execution until data is received)

The network model

- The processors of an asynchronous network coordinate their activities through **message passing**
 - A pair of processors need not be adjacent
 - **Routing** algorithms transmit a message through a network
- The topological properties of the network affects the system's processing capabilities:
 - **Diameter:** maximum distance between any two nodes
 - **Maximum degree:** of any node in G
 - **Node and edge connectivity:** the minimum number of nodes (edges) whose removal disconnects the graph
- We will briefly look at some representative topologies:
 - Linear array
 - 2D mesh
 - Hypercube


Linear array

- In a **linear array**, processor P_i is connected to P_{i-1} and P_{i+1} , if they exist
- In a **ring**, P_1 and P_p (the last processor) are connected to each other
- **Diameter:**
- **Maximum degree:**




Linear array

- In a **linear array**, processor P_i is connected to P_{i-1} and P_{i+1} , if they exist
- In a **ring**, P_1 and P_p (the last processor) are connected to each other
- **Diameter:**
 - $p - 1$
- **Maximum degree:**



Linear array

- In a **linear array**, processor P_i is connected to P_{i-1} and P_{i+1} , if they exist
- In a **ring**, P_1 and P_p (the last processor) are connected to each other
- **Diameter:**
 - $p - 1$
- **Maximum degree:**
 - 2



Linear array – matrix-vector multiplication

- Assume we want to compute $y = Ax$
 - With $p \leq n$
 - and $r = p/n$ an integer
- First, partition:
 - $A = (A_1, A_2, \dots, A_p)$
 - $x = (x_1, x_2, \dots, x_p)$
- Then, compute $z_i = A_i x_i$ on each processor independently
- Finally, add up the sum

$$z = \sum_{i=1}^p z_i$$

ALGORITHM 1.4

(Asynchronous Matrix Vector Product on a Ring)

Input: (1) The processor number i ; (2) the number p of processors; (3) the i th submatrix $B = A(1:n, (i-1)r+1:ir)$ of size $n \times r$, where $r = n/p$; (4) the i th subvector $w = x((i-1)r+1:ir)$ of size r .

Output: Processor P_i computes the vector $y = A_1 x_1 + \dots + A_i x_i$ and passes the result to the right. When the algorithm terminates, P_1 will hold the product Ax .

begin

1. Compute the matrix vector product $z = Bw$.
2. **if** $i = 1$ **then** set $y := 0$
 else **receive**(y , left)
3. Set $y := y + z$
4. **send**(y , right)
5. **if** $i = 1$ **then** **receive**(y , left)

end

Linear array – matrix-vector multiplication

We split the computations as follows (for $p = n/2$):

$$y_1 = a_{1,1}x_1 + a_{1,2}x_2 + a_{1,3}x_3 + a_{1,4}x_4 + \dots + a_{1,n-1}x_{n-1} + a_{1,n}x_n$$

$$y_2 = a_{2,1}x_1 + a_{2,2}x_2 + a_{2,3}x_3 + a_{2,4}x_4 + \dots + a_{2,n-1}x_{n-1} + a_{2,n}x_n$$

$$\vdots$$

$$y_n = a_{n,1}x_1 + a_{n,2}x_2 + a_{n,3}x_3 + a_{n,4}x_4 + \dots + a_{n,n-1}x_{n-1} + a_{n,n}x_n$$

for processors P_1, P_2, \dots, P_p resp.

Linear array – matrix-vector multiplication

We split the computations as follows:

$$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} z_{1,1} \\ z_{2,1} \\ \vdots \\ z_{n,1} \end{bmatrix} + \begin{bmatrix} z_{1,2} \\ z_{2,2} \\ \vdots \\ z_{n,2} \end{bmatrix} + \dots + \begin{bmatrix} z_{1,p} \\ z_{2,p} \\ \vdots \\ z_{n,p} \end{bmatrix}$$

for processors P_1, P_2, \dots, P_p resp.

Linear array – matrix-vector multiplication

- **Computation time:**

$$T_{comp} = O(n^2/p)$$

- Approx $\alpha(n^2/p)$ for some constant α

- However, P_1 has to wait until the $p - 1$ partial sums have been transmitted to execute the last instruction

- **Communication time:**

$$T_{comm} = p * comm(n)$$

- $comm(n)$ is the time needed to transmit n numbers between adjacent processors

ALGORITHM 1.4

(Asynchronous Matrix Vector Product on a Ring)

Input: (1) The processor number i ; (2) the number p of processors; (3) the i th submatrix $B = A(1 : n, (i - 1)r + 1 : ir)$ of size $n \times r$, where $r = n/p$; (4) the i th subvector $w = x((i - 1)r + 1 : ir)$ of size r .

Output: Processor P_i computes the vector $y = A_1x_1 + \dots + A_ix_i$ and passes the result to the right. When the algorithm terminates, P_1 will hold the product Ax .

begin

1. Compute the matrix vector product $z = Bw$.
2. **if** $i = 1$ **then** set $y := 0$
 else **receive**(y , left)
3. Set $y := y + z$
4. **send**(y , right)
5. **if** $i = 1$ **then** **receive**(y , left)

end

Linear array – matrix-vector multiplication

- $comm(n) \approx \sigma + n\tau$
 - σ : startup time
 - τ : transfer rate

- Total execution time:

$$T = T_{comp} + T_{comm}$$

$$\approx \alpha(n^2/p) + p(\sigma + n\tau)$$

- There is a trade-off between the two terms
- The sum is minimized when $\alpha(n^2/p) = p(\sigma + n\tau)$
 - Such that $p = n\sqrt{\alpha/(\sigma + n\tau)}$

ALGORITHM 1.4

(Asynchronous Matrix Vector Product on a Ring)

Input: (1) The processor number i ; (2) the number p of processors; (3) the i th submatrix $B = A(1:n, (i-1)r+1:ir)$ of size $n \times r$, where $r = n/p$; (4) the i th subvector $w = x((i-1)r+1:ir)$ of size r .

Output: Processor P_i computes the vector $y = A_1x_1 + \dots + A_ix_i$ and passes the result to the right. When the algorithm terminates, P_1 will hold the product Ax .

begin

1. Compute the matrix vector product $z = Bw$.
2. **if** $i = 1$ **then** set $y := 0$
 else **receive**(y , left)
3. Set $y := y + z$
4. **send**(y , right)
5. **if** $i = 1$ **then** **receive**(y , left)

end