

CSc 8530 Parallel Algorithms

Spring 2019

February 5th, 2019

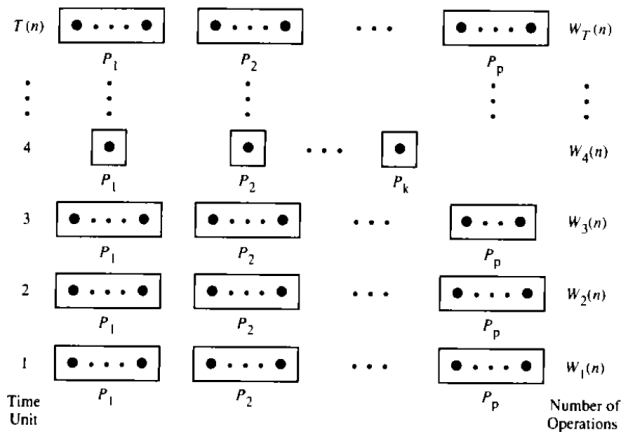
Worst-case analysis

- Let Q be a problem that we can solve in $T(n)$ with $P(n)$ processors
- **Parallel cost:** $C(n) = T(n)P(n)$
- The parallel algorithm can be converted to a sequential algorithm that runs in $O(C(n))$
- More generally, we can simulate a single step in $O(P(n)/p)$ sub-steps:
 - In sub-step 1: simulate processors $[1, p]$
 - In sub-step 2: simulate processors $[p + 1, 2p]$, etc.
- We can simulate the entire process in $O(T(n)P(n)/p)$

Work-time (WT) paradigm

- The **work-time (WT) paradigm** provides a two-level description of parallel algorithms
 - Upper level suppresses specific details
 - Lower level follows a general **scheduling principle**
- **Upper Level:** Describe the algorithm in terms of a sequence of time units
 - Each time unit may include any number of concurrent operations
- **Work:** total number of operations
- For convenience, at this level we can use a **pardo** statement
 - **for** $l \leq i \leq u$ **pardo** {statement(s)}
 - All the statements, for all valid indices, are executed concurrently

WT Scheduling Principle



WT vs. lower-level pseudocode

ALGORITHM 1.7

(Sum)

Input: $n = 2^k$ numbers stored in an array A .

Output: The sum $S = \sum_{i=1}^n A(i)$

begin

1. for $1 \leq i \leq n$ **parallel**
 Set $B(i) := A(i)$
2. for $h = 1$ to $\log n$ **do**
 for $1 \leq i \leq n/2^h$ **parallel**
 Set $B(i) := B(2i - 1) + B(2i)$
3. Set $S := B(1)$

end

WT pseudocode

ALGORITHM 1.8

(Sum Algorithm for Processor P_s)

Input: An array A of size $n = 2^k$ stored in the shared memory. The initialized local variables are (1) the order n ; (2) the number p of processors, where $p = 2^q \leq n$, and (3) the processor number s .

Output: The sum of the elements of A stored in the shared variable S . The array A retains its original value.

begin

1. for $j = 1$ to $l \left(\frac{n}{p} \right)$ **do**
 Set $B(l(s - 1) + j) := A(l(s - 1) + j)$
2. for $h = 1$ to $\log n$ **do**
 2.1. if $(k - h - q \geq 0)$ **then**
 for $j = 2^{k-h-q}(s - 1) + 1$ to $2^{k-h-q}s$ **do**
 Set $B(j) := B(2j - 1) + B(2j)$
- 2.2. **else** {if $(s \leq 2^{k-h})$ **then**
 Set $B(s) := B(2s - 1) + B(2s)$ }
3. if $(s = 1)$ **then** set $S := B(1)$

end

Lower-level pseudocode

Work vs. cost

- If a parallel algorithm runs in $T(n)$ with a total of $W(n)$ operations
 - Can be simulated in $O(\frac{W(n)}{p} + T(n))$ on a p -processor PRAM
 - The cost is $C_p(n) = T_p(n)p = O(W(n) + T(n)p)$
- Work and cost coincide asymptotically for $p = O(\frac{W(n)}{T(n)})$
- Otherwise they differ:
 - Work is independent of the number of processors
 - Cost is measured relative to the number of available processors
 - Cost \geq Work due to inefficient processor utilization
- For computing the sum of n numbers:
 - Work: $O(n)$, running time: $O(\log(n))$
 - Cost: $C_p(n) = O(n + p \log(n))$
 - With n processors, the cost is $O(n \log(n))$, not $O(n)$ (Why?)
 - We cannot use all the processors at all time steps, so the cost is higher than the total work

Optimality notions

- A sequential algorithm is **time optimal** iff its running time $T^*(n)$ cannot be improved asymptotically
- Two notions of optimality for parallel algorithms:
 - **Weak:** a WT presentation level algorithm is optimal iff $W(n) = \Theta(T^*(n))$
 - The total number of operations (not the running time) of the parallel algorithm is asymptotically equivalent to the sequential one
 - **Strong:** The running time $T(n)$ cannot be improved by any other parallel algorithm

Algorithmic techniques

- Designing parallel algorithms involves additional challenges compared to sequential methods
- We will now review some basic techniques for breaking down a problem into parallel chunks
- The example problems will often arise as sub-problems in more complicated applications

Balanced trees

- We have already encountered balanced binary trees
 - e.g., for summing the values of an array
- **General strategy:** *build a balanced binary tree on the input elements and traverse the tree forwards and backwards*
- An internal node u usually holds information about the data stored in the leaves of the subtree rooted at u
 - This strategy is useful when we can calculate this information quickly

Example: prefix sums

- Let $S = \{x_1, x_2, \dots, x_n\}$ be an n -element set
- Let $*$ be a binary associate operation (e.g., sum or product)
- A **prefix sum** is the partial sum defined by:

$$s_i = x_1 * x_2 * \dots * x_i, 1 \leq i \leq n$$

- The **prefix sums** are the n partial products s_1 to s_n
- A trivial sequential algorithm can compute s_i from s_{i-1} as
 $s_i = s_{i-1} * x_i$
 - Clearly, this algorithm is $O(n)$

Example: prefix sums

- We can use a balanced binary tree to compute the prefix sums in $O(\log(n))$
- We compute pairwise $*$ operations during the forward pass
- Each internal node will hold the sum of the elements stored in the leaves of its subtree
- During the backward pass, we compute the prefix sums at each level of the tree

Recursive prefix-sums algorithm

ALGORITHM 2.1

(Prefix Sums)

Input: An array of $n = 2^k$ elements (x_1, x_2, \dots, x_n) , where k is a nonnegative integer.

Output: The prefix sums s_i , for $1 \leq i \leq n$.

begin

1. **if** $n = 1$ **then** {set $s_1 := x_1$; **exit**}

2. **for** $1 \leq i \leq n/2$ **pardo**

 Set $y_i := x_{2i-1} * x_{2i}$

3. Recursively, compute the prefix sums of $\{y_1, y_2, \dots, y_{n/2}\}$, and store them in $z_1, z_2, \dots, z_{n/2}$.

4. **for** $1 \leq i \leq n$ **pardo**


i even : set $s_i := z_{i/2}$

i = 1 : set $s_1 := x_1$

i odd > 1 : set $s_i := z_{(i-1)/2} * x_i$


end

Recursive prefix-sums algorithm




Recursive prefix-sums algorithm

- **First time unit:** we compute $y_1 = x_1 * x_2$, etc.
- **Second:** $y'_1 = y_1 * y_2$ and $y'_2 = y_3 * y_4$
- **Third:** $y''_1 = y'_1 * y'_2$
- **Fourth:** We generate the prefix sum of the n elements



Recursive prefix-sums algorithm

- **Fifth:** generate z'_1 and z'_2 from y'_1 and y'_2 , resp.
- **Sixth:** z_1 to z_4 from y_1 to y_4
- **Seventh:** compute s_i for each i using the z and x values



Recursive prefix-sums algorithm – analysis

- For inputs of size $n = 2^k$, this algorithm requires $2^k + 1$ time units
 - We move up the k -level binary tree in the forward step
 - And down k steps in the backward step
- The algorithm can run in-place
 - In the sense that the y and z variables can be mapped to each other
- It is straightforward to show that this algorithm has:
 - **Running time:**
 - **Work:**

Recursive prefix-sums algorithm – analysis

- For inputs of size $n = 2^k$, this algorithm requires $2^k + 1$ time units
 - We move up the k -level binary tree in the forward step
 - And down k steps in the backward step
- The algorithm can run in-place
 - In the sense that the y and z variables can be mapped to each other
- It is straightforward to show that this algorithm has:
 - **Running time:**
 - $T(n) = O(\log(n))$
 - **Work:**

Recursive prefix-sums algorithm – analysis

- For inputs of size $n = 2^k$, this algorithm requires $2^k + 1$ time units
 - We move up the k -level binary tree in the forward step
 - And down k steps in the backward step
- The algorithm can run in-place
 - In the sense that the y and z variables can be mapped to each other
- It is straightforward to show that this algorithm has:
 - **Running time:**
 - $T(n) = O(\log(n))$
 - **Work:**
 - $W(n) = O(n)$

Recursive prefix-sums algorithm – analysis

- **Proof by induction:**
- The base case $k = 0$ is handled by step 1 of the algorithm
- Assume the algorithm works for $n = 2^k$
- We will prove it computes the prefix sums for $n = 2^{k+1}$
- The variables $z_1, z_2, \dots, z_{n/2}$ hold the prefix sums of the sequence $\{y_1, y_2, \dots, y_{n/2}\}$
- In particular,

$$\begin{aligned} z_j &= y_1 * y_2 * \dots * y_j \\ &= x_1 * x_2 * \dots * x_{2j-1} * x_{2j} \end{aligned}$$

- Thus, $z_j = s_{2j}$ for $1 \leq j \leq n/2$

Recursive prefix-sums algorithm – analysis

- **Proof by induction:**
- If i is even, then $s_i = z_{i/2}$
- If i is odd (and > 1):

$$\begin{aligned} s_i &= s_{2j+1} \\ &= s_{2j} * x_{2j+1} \\ &= z_{(i-1)/2} * x_i \end{aligned}$$

- All cases are handled appropriately, thus the algorithm works correctly for all inputs

Recursive prefix-sums algorithm – analysis

- **Resources required:**
- Step 1 takes $O(1)$ (sequential) time
- Steps 2 and 4 take $O(1)$ (parallel) time
 - With $O(n)$ operations per step
- Thus, the running time and work satisfy the following recurrences:

$$T(n) = T\left(\frac{n}{2}\right) + a$$

$$W(n) = W\left(\frac{n}{2}\right) + bn$$

where a and b are constants

- Their respective solutions are:

$$T(n) = O(\log n)$$

We reduce $T(n)$ by half in each step

$$W(n) = O(n)$$

The sum at each level decreases geometrically

Non-recursive prefix-sum algorithm

- The previous algorithm was recursive
- We can easily develop a non-recursive, yet still parallel version
- Here, we use auxiliary arrays $B(h, j)$ (forward values) and $C(h, j)$ (backward values) to simplify data storage
- Where $0 \leq h \leq \log(n)$ and $1 \leq j \leq n/2^h$
- For simplicity of analysis, we assume $n = 2^k$, for some k

Non-recursive prefix-sums algorithm

ALGORITHM 2.2

(Nonrecursive Prefix Sums)

Input: An array A of size $n = 2^k$, where k is a nonnegative integer.

Output: An array C such that $C(0, j)$ is the j th prefix sum, for $1 \leq j \leq n$.

begin

1. **for** $1 \leq j \leq n$ **pardo**

 Set $B(0, j) := A(j)$

2. **for** $h = 1$ **to** $\log n$ **do**

for $1 \leq j \leq n/2^h$ **pardo**

 Set $B(h, j) := B(h - 1, 2j - 1) * B(h - 1, 2j)$

3. **for** $h = \log n$ **to** 0 **do**


for $1 \leq j \leq n/2^h$ **pardo**

$\left\{ \begin{array}{l} j \text{ even} \quad : \text{Set } C(h, j) := C(h + 1, \frac{j}{2}) \\ j = 1 \quad : \text{Set } C(h, 1) := B(h, 1) \end{array} \right.$

$\left\{ \begin{array}{l} j \text{ odd} > 1 : \text{Set } C(h, j) := C(h + 1, \frac{j-1}{2}) * B(h, j) \end{array} \right.$


end

Non-recursive prefix-sums algorithm



- Similar to the problem of summing an individual array
- In the non-recursive version, we store all intermediate values in an auxiliary array

Non-recursive prefix-sums algorithm



- Here, we recurse forwards and backwards

Recursive vs. non-recursive versions

ALGORITHM 2.1

(Prefix Sums)

Input: An array of $n = 2^k$ elements (x_1, x_2, \dots, x_n) , where k is a nonnegative integer.

Output: The prefix sums s_i , for $1 \leq i \leq n$.

```
begin
1. if  $n = 1$  then {set  $s_1 := x_1$ ; exit}
2. for  $1 \leq i \leq n/2$  pardo
   Set  $y_i := x_{2i-1} * x_{2i}$ 
3. Recursively, compute the prefix sums of  $\{y_1, y_2, \dots, y_{n/2}\}$ , and
   store them in  $z_1, z_2, \dots, z_{n/2}$ .
4. for  $1 \leq i \leq n$  pardo
   {  $i$  even      : set  $s_i := z_{i/2}$ 
      $i = 1$       : set  $s_i := x_i$ 
      $i$  odd  $> 1$  : set  $s_i := z_{(i-1)/2} * x_i$  }
```

end

Recursive

ALGORITHM 2.2

(Nonrecursive Prefix Sums)

Input: An array A of size $n = 2^k$, where k is a nonnegative integer.

Output: An array C such that $C(0, j)$ is the j th prefix sum, for $1 \leq j \leq n$.

```
begin
1. for  $1 \leq j \leq n$  pardo
   Set  $B(0, j) := A(j)$ 
2. for  $h = 1$  to  $\log n$  do
   for  $1 \leq j \leq n/2^h$  pardo
     Set  $B(h, j) := B(h-1, 2j-1) * B(h-1, 2j)$ 
3. for  $h = \log n$  to 0 do
   for  $1 \leq j \leq n/2^h$  pardo
     {  $j$  even      : Set  $C(h, j) := C(h+1, \frac{j}{2})$ 
        $j = 1$       : Set  $C(h, 1) := B(h, 1)$ 
        $j$  odd  $> 1$  : Set  $C(h, j) := C(h+1, \frac{j-1}{2}) * B(h, j)$  }
```

end

Non-recursive

Review

- Building a balanced (binary) tree is a fundamental technique
- One of the most useful in parallel processing
- Other example problems:
 - Broadcasting a value to all processors
 - Compacting the labeled elements of an array
- Can be generalized to non-binary trees
 - Example: computing the maximum of n elements