

# CSc 8530

## Parallel Algorithms

Spring 2019

February 12th, 2019

# Root finding: parallel algorithm

## ALGORITHM 2.4

### (Pointer Jumping)

**Input:** A forest of rooted directed trees, each with a self-loop at its root, such that each arc is specified by  $(i, P(i))$ , where  $1 \leq i \leq n$ .


**Output:** For each vertex  $i$ , the root  $S(i)$  of the tree containing  $i$ .

**begin**


```
1. for  $1 \leq i \leq n$  pardo
    Set  $S(i) := P(i)$ 
    while  $(S(i)) \neq S(S(i))$  do
        Set  $S(i) := S(S(i))$ 
```

**end**

# Root finding: examples



(a)



- Notice how the distance to the root is cut in half in each iteration

# Root finding: analysis

- Let  $h$  be the maximum height of any tree in  $F$ 
  - In the worst case,  $h = O(n)$  (why?)
- Clearly, at the end  $S(i)$  is a root, for all  $i$ 
  - Intuitively, the roots are attractors or steady states of the iterative process
- The convergence rate to a root is  $2^h$ 
  - Because we half the distance to the root at each iteration
  - Hence, the number of iterations is  $O(\log(h))$
- Each iteration takes  $O(1)$  parallel time for  $O(n)$  operations (why?)
- **Running time:**
  - $O(\log(h))$
- **Work:**
  - $O(n \log(h))$

# Root finding

- We can easily show that this algorithm is not optimal
  - A  $T_s(n) = O(n)$  sequential algorithm exists
  - $W(n) > T_s(n)$ , so the algorithm cannot be **weakly** optimal
  - Intuitively, we are creating more work for ourselves by operating in parallel
- Could this algorithm be **strongly** optimal?
  - Theoretically yes, if  $T(n)$  cannot be improved by *any* other parallel algorithm

## Example: parallel prefix

- Assume that each node has a weight  $W(i)$ 
  - Unfortunately, the book uses  $W(i)$  for this quantity
  - Not to be confused with  $W(n)$ , which is work
- We can easily adapt the pointer jumping technique to compute the sum of weights from each node to its corresponding root
  - These sums are similar to the prefix sums we studied last class
- Intuitively, we update both the weights and the successors at the same time

# Parallel prefix using pointer jumping

## ALGORITHM 2.5

### (Parallel Prefix on Rooted Directed Trees)

**Input:** A forest of rooted directed trees, each with a self-loop at its root such that (1) each arc is specified by  $(i, P(i))$ , (2) each vertex  $i$  has a weight  $W(i)$ , and (3) for each root  $r$ ,  $W(r) = 0$ .

**Output:** For each vertex  $i$ ,  $W(i)$  is set equal to the sum of the weights of vertices on the path from  $i$  to the root of its tree.

**begin**

1. **for**  $1 \leq i \leq n$  **pardo**

    Set  $S(i) := P(i)$

**while**  $(S(i)) \neq S(S(i))$  **do**

        Set  $W(i) := W(i) + W(S(i))$

        Set  $S(i) := S(S(i))$

**end**

- The running time and work are asymptotically the same as the previous algorithm
  - Because we only added an  $O(1)$  operation

# Parallel divide and conquer


- Divide and conquer is one of the most natural ways to exploit parallelism
- Divide and conquer requires three steps:
  - 1 Partition the input into sub-partitions (ideally of similar size)
  - 2 Recursively solve the problem for each sub-partition
    - Remember that we can always replace recursion with stack-based iteration
  - 3 Combine the subproblem solutions into a solution for the overall problem
- Steps 1 and 2 are amenable for parallelization
- Step 3 will typically be more sequential



## Example: the convex hull problem

- Let  $S = \{p_1, p_2, \dots, p_n\}$  be  $n$  points in the plane
  - i.e.,  $p_i = (x_i, y_i)$ , for all  $i$
- The planar convex hull of  $S$  is the smallest convex polygon that contains all of  $S$
- Formally, a polygon  $Q$  is convex if the line segment connecting any two points inside it lies wholly within  $Q$
- The convex-hull problem is to determine the *ordered* list  $CH(S)$  of points that define the convex hull's boundary

# Convex hull example



- $CH(S) = \{v_1, v_2, \dots, v_8\}$

# The convex hull problem

- We can use a divide and conquer strategy to sequentially find  $CH(S)$  in  $O(n \log(n))$ 
  - Furthermore, sorting can be reduced to the convex hull problem
  - Thus  $T^*(n) = \Theta(n \log(n))$
- We can directly parallelize this sequential approach
  - Intuitively, we assign subsets of points to each processor
- The parallel algorithm is thus weakly optimal (why?)


# Parallel convex-hull algorithm

- Let  $p$  and  $q$  be the points in  $S$  with the largest and smallest  $x$  coordinates, resp.
- Clearly,  $p, q \in CH(S)$
- $p$  and  $q$  divide the convex hull into two parts:
  - **Upper hull**  $UH(S)$ : all the points from  $p$  to  $q$  (clockwise)
  - **Lower hull**  $LH(S)$ : all the points from  $q$  to  $p$  (clockwise)
- For simplicity, we will focus on finding  $UH(S)$ 
  - The problem for  $LH(S)$  is identical
- We also assume for simplicity that all  $x$  and  $y$  coordinates are unique and that  $n = 2^k$

# Parallel convex-hull algorithm

- We first sort all the points by their  $x$  coordinates
  - As noted, this takes  $\Theta(n \log(n))$
- We can parallel sort  $n$  numbers in  $T(n) = O(\log(n))$  and  $W(n) = O(n \log(n))$ 
  - We will prove this result in a few classes
- Let  $x(p_1) < x(p_2) < \dots < x(p_n)$
- Let  $S_1 = (p_1, p_2, \dots, p_{n/2})$  and  $S_2 = (p_{n/2+1}, \dots, p_n)$
- Suppose that we already have  $UH(S_1)$  and  $UH(S_2)$
- Then, the **upper common tangent** is the closest line that lies above both  $UH(S_1)$  and  $UH(S_2)$

# Upper common tangent example



- Both  $a$  and  $b$  have to be part of the convex hull of  $S$

# Parallel convex hull algorithm

- Determining the upper common tangent can be done sequentially with binary search ( $O(\log(n))$ )
  - We will explore faster parallel alternatives later in the course
- Let  $UH(S_1) = (q_1, \dots, q_s)$  and  $UH(S_2) = (q'_1, \dots, q'_t)$  be the left-to-right sorted lists of points
- Note that  $q_1 = p_1$  and  $q'_t = p_n$
- Assume that the upper common tangent is  $(q_i, q'_j)$
- Then,  $UH(S) = \{q_1, q_2, \dots, q_i, q'_j, \dots, q'_t\}$
- Once we find  $(q_i, q'_j)$ , we can compute  $UH(S)$  and its size in  $O(1)$  parallel time, using  $O(n)$  operations

# Parallel convex hull pseudocode

## ALGORITHM 2.6

### (Simple Upper Hull)

**Input:** A set  $S$  of  $n$  points in the plane, no two of which have the same  $x$  or  $y$  coordinates such that  $x(p_1) < x(p_2) < \dots < x(p_n)$ , where  $n$  is a power of 2.

**Output:** The upper hull of  $S$ .

**begin**

1. If  $n \leq 4$ , then use a brute-force method to determine  $UH(S)$ , and exit.
2. Let  $S_1 = (p_1, p_2, \dots, p_{\frac{n}{2}})$  and  $S_2 = (p_{\frac{n}{2}+1}, \dots, p_n)$ . Recursively, compute  $UH(S_1)$  and  $UH(S_2)$  in parallel.
3. Find the upper common tangent between  $UH(S_1)$  and  $UH(S_2)$ , and deduce the upper hull of  $S$ .

**end**

- The pseudocode for  $LH(S)$  is identical
- The final convex hull is then given by
$$CH(S) = UH(S) \cup LH(S)$$



# Parallel convex hull – analysis

- We can prove that the previous algorithm is correct by induction
  - The base case is valid by brute force
  - Then, intuitively, if  $UH(S_1)$  and  $UH(S_2)$  are correct, the upper common tangent method will connect them correctly
    - Because, by definition, every point lies on the lower half plane defined by the tangent
- Assume the algorithm takes  $T(n)$  steps using  $W(n)$  work
  - We will now define these values in terms of a recurrence relation

# Parallel convex hull – analysis

- Step 1 takes  $O(1)$  sequential time
- Step 2 takes  $T(n/2)$  using  $2W(n/2)$  operations
- Step 3:
  - Upper common tangent:  $O(\log(n))$
  - Combining upper hulls:  $O(1)$  parallel time and  $O(n)$  work
- Thus:

$$T(n) \leq T\left(\frac{n}{2}\right) + a \log(n)$$

$$W(n) \leq 2W\left(\frac{n}{2}\right) + bn$$

$a$  and  $b$  are positive constants

- Total running time and work:

# Parallel convex hull – analysis

- Step 1 takes  $O(1)$  sequential time
- Step 2 takes  $T(n/2)$  using  $2W(n/2)$  operations
- Step 3:
  - Upper common tangent:  $O(\log(n))$
  - Combining upper hulls:  $O(1)$  parallel time and  $O(n)$  work
- Thus:

$$T(n) \leq T\left(\frac{n}{2}\right) + a \log(n)$$

$$W(n) \leq 2W\left(\frac{n}{2}\right) + bn$$

$a$  and  $b$  are positive constants

- Total running time and work:
  - $T(n) = O(\log^2(n))$

# Parallel convex hull – analysis

- Step 1 takes  $O(1)$  sequential time
- Step 2 takes  $T(n/2)$  using  $2W(n/2)$  operations
- Step 3:
  - Upper common tangent:  $O(\log(n))$
  - Combining upper hulls:  $O(1)$  parallel time and  $O(n)$  work
- Thus:

$$T(n) \leq T\left(\frac{n}{2}\right) + a \log(n)$$

$$W(n) \leq 2W\left(\frac{n}{2}\right) + bn$$

$a$  and  $b$  are positive constants

- Total running time and work:
  - $T(n) = O(\log^2(n))$
  - $W(n) = O(n \log(n))$

## Parallel convex hull – other analyses

- The previous algorithm requires the CREW PRAM model (why?)
- For  $p$  processors, the algorithm runs in  $O\left(\frac{n \log(n)}{p} + \log^2(n)\right)$ 
  - What values of  $p$  achieve an optimal speedup?

## Parallel convex hull – other analyses

- The previous algorithm requires the CREW PRAM model (why?)
  - Merging  $UH(S_1)$  and  $UH(S_2)$  may require accessing the same point simultaneously
- For  $p$  processors, the algorithm runs in  $O\left(\frac{n \log(n)}{p} + \log^2(n)\right)$ 
  - What values of  $p$  achieve an optimal speedup?

# Parallel convex hull – other analyses

- The previous algorithm requires the CREW PRAM model (why?)
  - Merging  $UH(S_1)$  and  $UH(S_2)$  may require accessing the same point simultaneously
- For  $p$  processors, the algorithm runs in  $O\left(\frac{n \log(n)}{p} + \log^2(n)\right)$ 
  - What values of  $p$  achieve an optimal speedup?
  - We want  $\frac{n \log(n)}{p} \leq \log^2(n)$  to eliminate the left-hand term

# Parallel convex hull – other analyses

- The previous algorithm requires the CREW PRAM model (why?)
  - Merging  $UH(S_1)$  and  $UH(S_2)$  may require accessing the same point simultaneously
- For  $p$  processors, the algorithm runs in  $O\left(\frac{n \log(n)}{p} + \log^2(n)\right)$ 
  - What values of  $p$  achieve an optimal speedup?
  - We want  $\frac{n \log(n)}{p} \leq \log^2(n)$  to eliminate the left-hand term
  - Thus:

$$\frac{n \log(n)}{p} \leq \log^2(n)$$

$$n \log(n) \leq p \log^2(n)$$

$$\frac{n}{\log(n)} \leq p$$

- Any value of  $p$  in this range will be asymptotically dominated by the  $\log^2(n)$  term