

# CSc 8530

## Parallel Algorithms

Spring 2019

February 14th, 2019

# Parallel prefix using pointer jumping

## ALGORITHM 2.5

### (Parallel Prefix on Rooted Directed Trees)

**Input:** A forest of rooted directed trees, each with a self-loop at its root such that (1) each arc is specified by  $(i, P(i))$ , (2) each vertex  $i$  has a weight  $W(i)$ , and (3) for each root  $r$ ,  $W(r) = 0$ .

**Output:** For each vertex  $i$ ,  $W(i)$  is set equal to the sum of the weights of vertices on the path from  $i$  to the root of its tree.

**begin**

1. **for**  $1 \leq i \leq n$  **pardo**

    Set  $S(i) := P(i)$

**while**  $(S(i)) \neq S(S(i))$  **do**

        Set  $W(i) := W(i) + W(S(i))$

        Set  $S(i) := S(S(i))$


**end**

- The running time and work are asymptotically the same as the previous algorithm
  - Because we only added an  $O(1)$  operation

# Parallel convex-hull algorithm

- We first sort all the points by their  $x$  coordinates
  - As noted, this takes  $\Theta(n \log(n))$
- We can parallel sort  $n$  numbers in  $T(n) = O(\log(n))$  and  $W(n) = O(n \log(n))$ 
  - We will prove this result in a few classes
- Let  $x(p_1) < x(p_2) < \dots < x(p_n)$
- Let  $S_1 = (p_1, p_2, \dots, p_{n/2})$  and  $S_2 = (p_{n/2+1}, \dots, p_n)$
- Suppose that we already have  $UH(S_1)$  and  $UH(S_2)$
- Then, the **upper common tangent** is the closest line that lies above both  $UH(S_1)$  and  $UH(S_2)$

# Upper common tangent example



- Both  $a$  and  $b$  have to be part of the convex hull of  $S$

# Parallel convex hull pseudocode

## ALGORITHM 2.6

### (Simple Upper Hull)

**Input:** A set  $S$  of  $n$  points in the plane, no two of which have the same  $x$  or  $y$  coordinates such that  $x(p_1) < x(p_2) < \dots < x(p_n)$ , where  $n$  is a power of 2.

**Output:** The upper hull of  $S$ .

**begin**

1. If  $n \leq 4$ , then use a brute-force method to determine  $UH(S)$ , and exit.
2. Let  $S_1 = (p_1, p_2, \dots, p_{\frac{n}{2}})$  and  $S_2 = (p_{\frac{n}{2}+1}, \dots, p_n)$ . Recursively, compute  $UH(S_1)$  and  $UH(S_2)$  in parallel.
3. Find the upper common tangent between  $UH(S_1)$  and  $UH(S_2)$ , and deduce the upper hull of  $S$ .

**end**

- The pseudocode for  $LH(S)$  is identical
- The final convex hull is then given by
$$CH(S) = UH(S) \cup LH(S)$$

# Parallel convex hull – analysis

- Step 1 takes  $O(1)$  sequential time
- Step 2 takes  $T(n/2)$  using  $2W(n/2)$  operations
- Step 3:
  - Upper common tangent:  $O(\log(n))$
  - Combining upper hulls:  $O(1)$  parallel time and  $O(n)$  work
- Thus:

$$T(n) \leq T\left(\frac{n}{2}\right) + a \log(n)$$

$$W(n) \leq 2W\left(\frac{n}{2}\right) + bn$$

$a$  and  $b$  are positive constants

- Total running time and work:
  - $T(n) = O(\log^2(n))$
  - $W(n) = O(n \log(n))$

# Parallel convex hull – other analyses

- The previous algorithm requires the CREW PRAM model (why?)
  - Merging  $UH(S_1)$  and  $UH(S_2)$  may require accessing the same point simultaneously
- For  $p$  processors, the algorithm runs in  $O\left(\frac{n \log(n)}{p} + \log^2(n)\right)$ 
  - What values of  $p$  achieve an optimal speedup?
  - We want  $\frac{n \log(n)}{p} \leq \log^2(n)$  to eliminate the left-hand term
  - Thus:

$$\frac{n \log(n)}{p} \leq \log^2(n)$$

$$n \log(n) \leq p \log^2(n)$$

$$\frac{n}{\log(n)} \leq p$$

- Any value of  $p$  in this range will be asymptotically dominated by the  $\log^2(n)$  term

# Upper common tangent – sequential algorithm


- We previously noted that can compute the upper common tangent sequentially in  $O(\log(n))$
- The actual  $O(\log(n))$  algorithms [Overmars & van Leeuwen, 1981] and [Kirkpatrick & Snoeyink, 1995] are rather technical
- Here, we will discuss an  $O(n)$  algorithm instead, which is much easier to follow
- Later in the course, we will study a constant time parallel version



# Upper common tangent – sequential algorithm

- First, let  $a$  and  $b$  be the rightmost and leftmost vertices of  $UH(S_1)$  and  $UH(S_2)$ , resp.
  - i.e.,  $a = UH(S_1)[i]$  and  $b = UH(S_2)[j]$ , for some  $i, j$
- Then, **while**  $(a, b)$  dips below either hull
  - **while**  $(a, b)$  dips below  $UH(S_2)$  (i.e., not tangent)
    - $b = UH(S_2)[j + 1]$
  - **while**  $(a, b)$  dips below  $UH(S_1)$  (i.e., not tangent)
    - $a = UH(S_1)[i - 1]$
- Intuitively, we alternate between moving  $b$  to the right and  $a$  to the left

# Upper common tangent – sequential algorithm



- Example taken from [O'Rourke, 1997]
- Here, we compute the lower common tangent, but the idea is the same

# Upper common tangent – sequential algorithm

- We can check if a point is above or below a line  $(a, b)$  by using the equation for a line:

$$\begin{aligned}(y - y_a)/(x - x_a) &= (y_b - y_a)/(x_b - x_a) \\ (y_b - y_a)x - (x_b - x_a)y &= x_a * y_b - x_b * y_a \\ \alpha x + \beta y &= \gamma\end{aligned}$$

- All the points such that  $\alpha x + \beta y < \gamma$  are below the line (and vice versa for above)
- For the UH and LH, we only need to check the two neighbors of the two current candidate points
  - Because the hulls are convex
  - Takes constant sequential time

# Partitioning strategy

- The **partitioning strategy** consists of:
  - 1 Breaking up a problem into  $p$  *independent* problems of roughly equal size
  - 2 Solving the subproblems concurrently
- Differs from divide and conquer:
  - The splits are not (necessarily) recursive
  - The main work lies in partitioning the input, not in combining the solutions of the subproblems
- In the simplest case, we simply break up the data into  $p$  non-overlapping chunks
- More generally, we ensure that the subproblems are independent, even if some of the data they access is the same

# Partitioning example – merging sorted sequences

- Given a set  $S$  with a partial order relation  $\leq$ ,  $S$  is **totally ordered** if,  $a \leq b$  or  $b \leq a$ , for all  $a, b \in S$
- Given two sorted sequences  $A$  and  $B$  drawn from  $S$ , we want to merge these two sequences into one
- The basic sequential algorithm is  $O(n)$ 
  - Repeatedly move the smallest of the two lists to the final list
- We will explore a parallel solution that partitions  $A$  and  $B$  into many pairs of subsequences

# A simple merging algorithm

- Let  $X = (x_1, x_2, \dots, x_t)$  be a sequence of elements drawn from  $S$  (not necessarily sorted)
- The rank of a new element  $y \notin X$ ,  $\text{rank}(y : X)$  is the number of elements of  $X$  that are less than or equal to  $y$
- Let  $Y = (y_1, y_2, \dots, y_s)$  be another list of elements drawn from  $S$
- We want to determine  $\text{rank}(Y : X)$ 
  - Example:  $X = (25, -13, 26, 31, 54, 7)$  and  $Y = (13, 27, -27)$

# A simple merging algorithm

- Let  $X = (x_1, x_2, \dots, x_t)$  be a sequence of elements drawn from  $S$  (not necessarily sorted)
- The rank of a new element  $y \notin X$ ,  $\text{rank}(y : X)$  is the number of elements of  $X$  that are less than or equal to  $y$
- Let  $Y = (y_1, y_2, \dots, y_s)$  be another list of elements drawn from  $S$
- We want to determine  $\text{rank}(Y : X)$ 
  - Example:  $X = (25, -13, 26, 31, 54, 7)$  and  $Y = (13, 27, -27)$
  - $\text{rank}(Y : X) = (2, 4, 0)$
  - Note,  $\text{rank}(y_i : X)$  is calculated independently for every element of  $Y$
  - The elements of  $Y$  are not inserted into  $X$

# A simple merging algorithm

- Here, we assume for simplicity that all elements of  $A$  and  $B$  are distinct
- The merging problem is equivalent to determining the rank of every element from  $A$  or  $B$  in the union  $A \cup B$ 
  - If  $\text{rank}(x : A \cup B) = i$ , then  $x$  should be placed in the  $i$  element of the combined list
- Thus,  $\text{rank}(x : A \cup B) = \text{rank}(x : A) + \text{rank}(x : B)$  (why?)
- We can solve the merging problem by determining  $\text{rank}(A : B)$  and  $\text{rank}(B : A)$  independently



# A simple merging algorithm

- Let  $b_i$  be an arbitrary element of  $B$
- Since  $A$  is sorted, we can find  $\text{rank}(b_i : A)$  using binary search
  - Runs in  $O(\log(n))$  (why?)
- If we run  $O(n)$  binary searches in parallel, we can solve the merge problem in:

# A simple merging algorithm

- Let  $b_i$  be an arbitrary element of  $B$
- Since  $A$  is sorted, we can find  $\text{rank}(b_i : A)$  using binary search
  - Runs in  $O(\log(n))$  (why?)
- If we run  $O(n)$  binary searches in parallel, we can solve the merge problem in:
  - $T(n) = O(\log(n))$
  - $W(n) = O(n \log(n))$
- The work is non-optimal (why?)


# A simple merging algorithm

- Let  $b_i$  be an arbitrary element of  $B$
- Since  $A$  is sorted, we can find  $\text{rank}(b_i : A)$  using binary search
  - Runs in  $O(\log(n))$  (why?)
- If we run  $O(n)$  binary searches in parallel, we can solve the merge problem in:
  - $T(n) = O(\log(n))$
  - $W(n) = O(n \log(n))$
- The work is non-optimal (why?)
  - The sequential algorithm is  $O(n)$

# An optimal merging algorithm

- We can design an optimal merging algorithm as follows:
  - ① Choose approximately  $n/\log(n)$  elements of each of  $A$  and  $B$  that partition  $A$  and  $B$  into blocks of almost equal lengths
  - ② Apply the binary search method to rank each of the chosen elements in the other sequence
- We reduce the problem to merging pairs of  $O(\log(n))$  sequences
- For simplicity, though, we will discuss a slight variant in which we only partition  $B$  into equal-sized blocks
  - The blocks of  $A$  may vary in size

# An optimal merging algorithm – partitioning illustration



- Each  $B_i$  is of size  $\log(m)$
- The  $A_j$  blocks could be of different sizes
- Here,  $j(i) = \text{rank}(b_{i \log(m)} : A)$ 
  - That is,  $A(j) \leq b_{i \log(m)}$ , for all  $j \leq j(i)$

# An optimal merging algorithm – partitioning pseudocode

## ALGORITHM 2.7

### (Partition)

**Input:** Two arrays  $A = (a_1, \dots, a_n)$  and  $B = (b_1, \dots, b_m)$  in increasing order, where both  $\log m$  and  $k(m) = m/\log m$  are integers.

**Output:**  $k(m)$  pairs  $(A_i, B_i)$  of subsequences of  $A$  and  $B$  such that (1)  $|B_i| = \log m$ , (2)  $\sum_i |A_i| = n$ , and (3) each element of  $A_i$  and  $B_i$  is larger than each element of  $A_{i-1}$  or  $B_{i-1}$ , for all  $1 \leq i \leq k(m) - 1$ .

**begin**

1. Set  $j(0) := 0, j(k(m)) := n$
2. **for**  $1 \leq i \leq k(m) - 1$  **pardo**
  - 2.1. Rank  $b_{i \log m}$  in  $A$  using the binary search method, and let  $j(i) = \text{rank}(b_{i \log m} : A)$
3. **for**  $0 \leq i \leq k(m) - 1$  **pardo**
  - 3.1. Set  $B_i := (b_{i \log m + 1}, \dots, b_{(i+1) \log m})$
  - 3.2. Set  $A_i := (a_{j(i)+1}, \dots, a_{j(i+1)})$   
( $A_i$  is empty if  $j(i) = j(i+1)$ )

**end**