

CSc 8530

Parallel Algorithms

Spring 2019

March 26th, 2019

Paper presentation guidelines

- The presentation will be graded based on the following criteria:
 - ① Content mastery (40%): how well your group understands the paper's content. Based on both your presentation and how you answer questions from the audience.
 - ② Preparedness (30%): slide quality; time management; smooth transitions between speakers, etc.
 - ③ Delivery (30%): quality of the oral presentation; effective use of visual aids, etc.
- **Note:** each person's grade will be **individual**.
- **Tips:**
 - ① As a rule of thumb, allocate around 60-90 seconds per slide
 - ② Aim for fewer words and more graphics in your slides
 - ③ Don't try to go over every detail in the paper (there is not enough time); focus on the big picture

Matrix multiplication efficiency

```
__global__ void MatrixMulKernel(float* M, float* N, float* P,
int Width) {
    // Calculate the row index of the P element and M
    int Row = blockIdx.y*blockDim.y+threadIdx.y;
    // Calculate the column index of P and N
    int Col = blockIdx.x*blockDim.x+threadIdx.x;
    if ((Row < Width) && (Col < Width)) {
        float Pvalue = 0;
        // each thread computes one element of the block sub-matrix
        for (int k = 0; k < Width; ++k) {
            Pvalue += M[Row*Width+k]*N[k*Width+Col];
        }
        P[Row*Width+Col] = Pvalue;
    }
}
```

- The memory accesses of this program are dominated by the inner for-loop
 - Two global memory accesses
 - One addition and one multiplication
- Compute ratio of 1.0


CUDA memory types

Variable declaration	Memory	Scope	Lifetime
Automatic variables other than arrays	Register	Thread	Kernel
Automatic array variables	Local	Thread	Kernel
<code>__device__ __shared__ int SharedVar;</code>	Shared	Block	Kernel
<code>__device__ int GlobalVar;</code>	Global	Grid	Application
<code>__device__ __constant__ int ConstVar;</code>	Constant	Grid	Application

- In CUDA, we can declare variables so that reside in a specific type of memory
- **Scope** is the set of threads that have access to it
- **Lifetime** is how long the variable is maintained
 - e.g., the top three are initialized every time we call a kernel function

Matrix multiplication revisited

- The colors indicate which threads need which parts of the input matrices
- The black arrows show the memory accesses of the first two threads, $P_{0,0}$ and $P_{0,1}$, from the first block
- They both need the same row, but different columns
- They access each needed element sequentially




Matrix multiplication revisited

Access order →

thread _{0,0}	$M_{0,0} * N_{0,0}$	$M_{0,1} * N_{1,0}$	$M_{0,2} * N_{2,0}$	$M_{0,3} * N_{3,0}$
thread _{0,1}	$M_{0,0} * N_{0,1}$	$M_{0,1} * N_{1,1}$	$M_{0,2} * N_{2,1}$	$M_{0,3} * N_{3,1}$
thread _{1,0}	$M_{1,0} * N_{0,0}$	$M_{1,1} * N_{1,0}$	$M_{1,2} * N_{2,0}$	$M_{1,3} * N_{3,0}$
thread _{1,1}	$M_{1,0} * N_{0,1}$	$M_{1,1} * N_{1,1}$	$M_{1,2} * N_{2,1}$	$M_{1,3} * N_{3,1}$


- Above, we have the sequence of memory accesses for all threads, T , in block(0,0)
- e.g. both $T_{0,0}$ and $T_{0,1}$ access $M_{0,0}$
- Note how, in this example each element of M and N is accessed twice
 - If we only retrieve each element once, we reduce global memory accesses by half

Matrix multiplication revisited



- Above, we divide M and N into 2×2 tiles
- We load one tile of M and one of N into the shared memory at a time
- In the simplest case, tile size equals block size
 - But it doesn't have to

Matrix multiplication revisited



- All the threads collaborate to load individual elements from the current tile
- Once elements are loaded onto shared memory, they can be accessed by multiple threads

Matrix multiplication revisited

	Phase 1			Phase 2		
thread _{0,0}	$M_{0,0}$ ↓ $Mds_{0,0}$	$N_{0,0}$ ↓ $Nds_{0,0}$	$PValue_{0,0} +=$ $Mds_{0,0} * Nds_{0,0} +$ $Mds_{0,1} * Nds_{1,0}$	$M_{0,2}$ ↓ $Mds_{0,0}$	$N_{2,0}$ ↓ $Nds_{0,0}$	$PValue_{0,0} +=$ $Mds_{0,0} * Nds_{0,0} +$ $Mds_{0,1} * Nds_{1,0}$
thread _{0,1}	$M_{0,1}$ ↓ $Mds_{0,1}$	$N_{0,1}$ ↓ $Nds_{1,0}$	$PValue_{0,1} +=$ $Mds_{0,0} * Nds_{0,1} +$ $Mds_{0,1} * Nds_{1,1}$	$M_{0,3}$ ↓ $Mds_{0,1}$	$N_{2,1}$ ↓ $Nds_{0,1}$	$PValue_{0,1} +=$ $Mds_{0,0} * Nds_{0,1} +$ $Mds_{0,1} * Nds_{1,1}$
thread _{1,0}	$M_{1,0}$ ↓ $Mds_{1,0}$	$N_{1,0}$ ↓ $Nds_{1,0}$	$PValue_{1,0} +=$ $Mds_{1,0} * Nds_{0,0} +$ $Mds_{1,1} * Nds_{1,0}$	$M_{1,2}$ ↓ $Mds_{1,0}$	$N_{3,0}$ ↓ $Nds_{1,0}$	$PValue_{1,0} +=$ $Mds_{1,0} * Nds_{0,0} +$ $Mds_{1,1} * Nds_{1,0}$
thread _{1,1}	$M_{1,1}$ ↓ $Mds_{1,1}$	$N_{1,1}$ ↓ $Nds_{1,1}$	$PValue_{1,1} +=$ $Mds_{1,0} * Nds_{0,1} +$ $Mds_{1,1} * Nds_{1,1}$	$M_{1,3}$ ↓ $Mds_{1,1}$	$N_{3,1}$ ↓ $Nds_{1,1}$	$PValue_{1,1} +=$ $Mds_{1,0} * Nds_{0,1} +$ $Mds_{1,1} * Nds_{1,1}$

time →

- We now have to compute the dot product in **phases**
 - We accumulate the partial dot product after every phase
- We progressively load parts of the rows and columns into smaller arrays Mds and Nds
 - These arrays live in shared memory
 - We can reuse them every time a new tile is loaded
 - Memory accesses exhibit **locality**

Matrix multiplication revisited

```

__global__ void MatrixMulKernel(float* d_M, float* d_N, float* d_P,
int Width) {

1.  __shared__ float Mds[TILE_WIDTH][TILE_WIDTH];
2.  __shared__ float Nds[TILE_WIDTH][TILE_WIDTH];

3.  int bx = blockIdx.x; int by = blockIdx.y;
4.  int tx = threadIdx.x; int ty = threadIdx.y;

    // Identify the row and column of the d_P element to work on
5.  int Row = by * TILE_WIDTH + ty;
6.  int Col = bx * TILE_WIDTH + tx;

7.  float Pvalue = 0;
    // Loop over the d_M and d_N tiles required to compute d_P element
8.  for (int ph = 0; ph < Width/TILE_WIDTH; ++ph) {

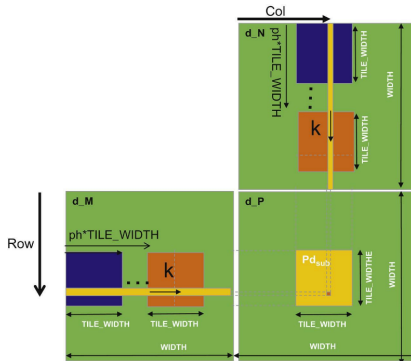
        // Collaborative loading of d_M and d_N tiles into shared memory
9.  Mds[ty][tx] = d_M[Row*Width + ph*TILE_WIDTH + tx];
10. Nds[ty][tx] = d_N[(ph*TILE_WIDTH + ty)*Width + Col];
11. __syncthreads();

12. for (int k = 0; k < TILE_WIDTH; ++k) {
13.     Pvalue += Mds[ty][k] * Nds[k][tx];
14. }
15. __syncthreads();
    }
    d_P[Row*Width + Col] = Pvalue;
}

```

- `__shared__` arrays `Mds` and `Nds` are common to the *block*
- `__syncthreads()`; ensures that all the elements in a tile have been loaded before the next phase of the dot product

Matrix multiplication revisited



- Our memory accesses are reduced by a factor of $TILE_WIDTH$
- For 16×16 tiles, this yields a compute ratio of 16
 - e.g., with 150 GB/s bandwidth, we can achieve $(150/4)16 = 600$ GFLOPS

Prefix sums revisited

- We will now revisit the prefix sums problem
 - In the context of GPUs
- As before, this is a *model* problem
 - Also called a *pattern* in software development
 - The techniques used to solve it are applicable to a wide variety of other problems
- We will investigate three different types of kernel functions:
 - Kogge-Stone
 - Brent-Kung
 - two-phase hybrid
- Each involves different computational tradeoffs

Prefix sums

- Let $S = \{x_1, x_2, \dots, x_n\}$ be an n -element set
- Let $*$ be a binary associate operation (e.g., sum or product)
- A **prefix sum** is the partial sum defined by:


$$s_i = x_1 * x_2 * \dots * x_i, 1 \leq i \leq n$$

- The **prefix sums** are the n partial products s_1 to s_n
- A trivial sequential algorithm can compute s_i from s_{i-1} as
$$s_i = s_{i-1} * x_i$$
 - Clearly, this algorithm is $O(n)$

Prefix sums

- We can use a balanced binary tree to compute the prefix sums in $O(\log(n))$
- We compute pairwise $*$ operations during the forward pass
- Each internal node will hold the sum of the elements stored in the leaves of its subtree
- During the backward pass, we compute the prefix sums at each level of the tree

Balanced binary trees



Recursive vs. non-recursive versions

ALGORITHM 2.1

(Prefix Sums)

Input: An array of $n = 2^k$ elements (x_1, x_2, \dots, x_n) , where k is a nonnegative integer.

Output: The prefix sums s_i , for $1 \leq i \leq n$.

begin

1. **if** $n = 1$ **then** $\{set\ s_1 := x_1; \text{exit}\}$
2. **for** $1 \leq i \leq n/2$ **pardo**
 $\quad Set\ y_i := x_{2i-1} * x_{2i}$
3. Recursively, compute the prefix sums of $\{y_1, y_2, \dots, y_{n/2}\}$, and store them in $z_1, z_2, \dots, z_{n/2}$.
4. **for** $1 \leq i \leq n$ **pardo**
 $\quad \{i\ \text{even} \quad : set\ s_i := z_{i/2}$
 $\quad \quad i = 1 \quad : set\ s_1 := x_1$
 $\quad i\ \text{odd} > 1 : set\ s_i := z_{(i-1)/2} * x_i\}$

end

Recursive

ALGORITHM 2.2

(Nonrecursive Prefix Sums)

Input: An array A of size $n = 2^k$, where k is a nonnegative integer.

Output: An array C such that $C(0, j)$ is the j th prefix sum, for $1 \leq j \leq n$.

begin

1. **for** $1 \leq j \leq n$ **pardo**
 $\quad Set\ B(0, j) := A(j)$
2. **for** $h = 1$ **to** $\log n$ **do**
 $\quad \text{for } 1 \leq j \leq n/2^h$ **pardo**
 $\quad \quad Set\ B(h, j) := B(h-1, 2j-1) * B(h-1, 2j)$
3. **for** $h = \log n$ **to** 0 **do**
 $\quad \text{for } 1 \leq j \leq n/2^h$ **pardo**
 $\quad \quad \{j\ \text{even} \quad : Set\ C(h, j) := C(h+1, \frac{j}{2})$
 $\quad \quad \quad j = 1 \quad : Set\ C(h, 1) := B(h, 1)$
 $\quad \quad j\ \text{odd} > 1 : Set\ C(h, j) := C(h+1, \frac{j-1}{2}) * B(h, j)\}$

end

Non-recursive

Parallel scans

- Parallel scans (i.e, a more general version of prefix sums) are often used to convert seemingly sequential operations into parallel ones
- Most recursive functions can be formulated as parallel scans
- In general, given an input set $S = \{x_1, x_2, \dots, x_n\}$ and a binary operation \oplus our output set is:

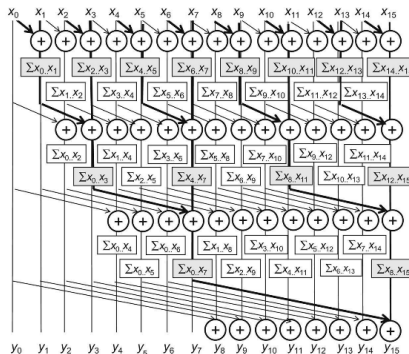
$$\{ \{x_1\}, \{x_1 \oplus x_2\}, \dots, \{x_1 \oplus x_2 \oplus \dots x_{n_1} \oplus x_n\} \}$$

Kogge-Stone – a simple parallel scan

- As noted before, we can compute a parallel scan by *reducing* the input elements
 - That is, merging partial results using a binary tree
- The Kogge-Stone algorithm is one of the simplest ways to build this tree
- At iteration 0, we assume position $X[i]$ in our input array contains element x_i
- At iteration n , $X[i]$ will contain the sum of 2^n elements leading to i (including x_i):
 - e.g., for $n = 2$, we have:

$$X[i] = x_{i-3} + x_{i-2} + x_{i-1} + x_i$$

Kogge-Stone – illustration



- The above illustrates the algorithm for a 16-element array
- At iteration j , each element adds its current value with the value of the element that is 2^j steps before it
 - If this element is out of bounds, then we stop computing values for that position

Kogge-Stone – code

```
__global__ void Kogge-Stone_scan_kernel(float *X, float *Y,
int InputSize) {

    __shared__ float XY[SECTION_SIZE];

    int i = blockIdx.x*blockDim.x + threadIdx.x;
    if (i < InputSize) {
        XY[threadIdx.x] = X[i];
    }

    // the code below performs iterative scan on XY
    for (unsigned int stride = 1; stride < blockDim.x; stride *= 2) {
        __syncthreads();
        if (threadIdx.x >= stride) XY[threadIdx.x] += XY[threadIdx.x-stride];
    }

    Y[i] = XY[threadIdx.x];
}
```

- The above code computes Kogge-Stone for a **section** of the array that is small enough to fit in a block
 - Each thread is responsible for one element of the output array
- We will see how to combine multiple blocks later

Kogge-Stone – speed and work efficiency

- We will now analyze the performance of the previous kernel
- All threads execute for $\log(n)$ steps, where n is `SECTION_SIZE`
 - Why not array size?
- In each iteration j , the number of *inactive* threads is equal to the stride size, 2^j
- Thus, the total work is:

$$\begin{aligned} W &= \sum_{j=0}^{\log(n)} n - 2^j \\ &= n \log(n) - (n - 1) \\ &= O(n \log(n)) \end{aligned}$$

- Compare to the sequential algorithm: $O(n)$